Object Categorization in the Sink:
Learning Behavior–Grounded Object Categories with Water

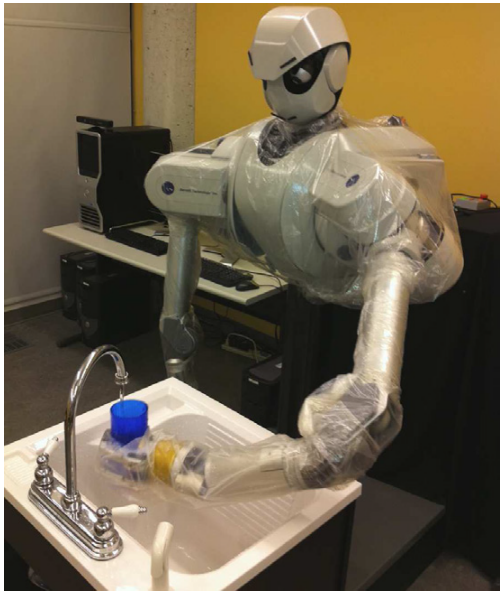

Shane Griffith, Vlad Sukhoy, Todd Wegter, and Alex Stoytchev

Developmental Robotics Laboratory
Iowa State University

www.ece.iastate.edu/~shaneg

Humanoid Robots and Water Don't Play Well Together

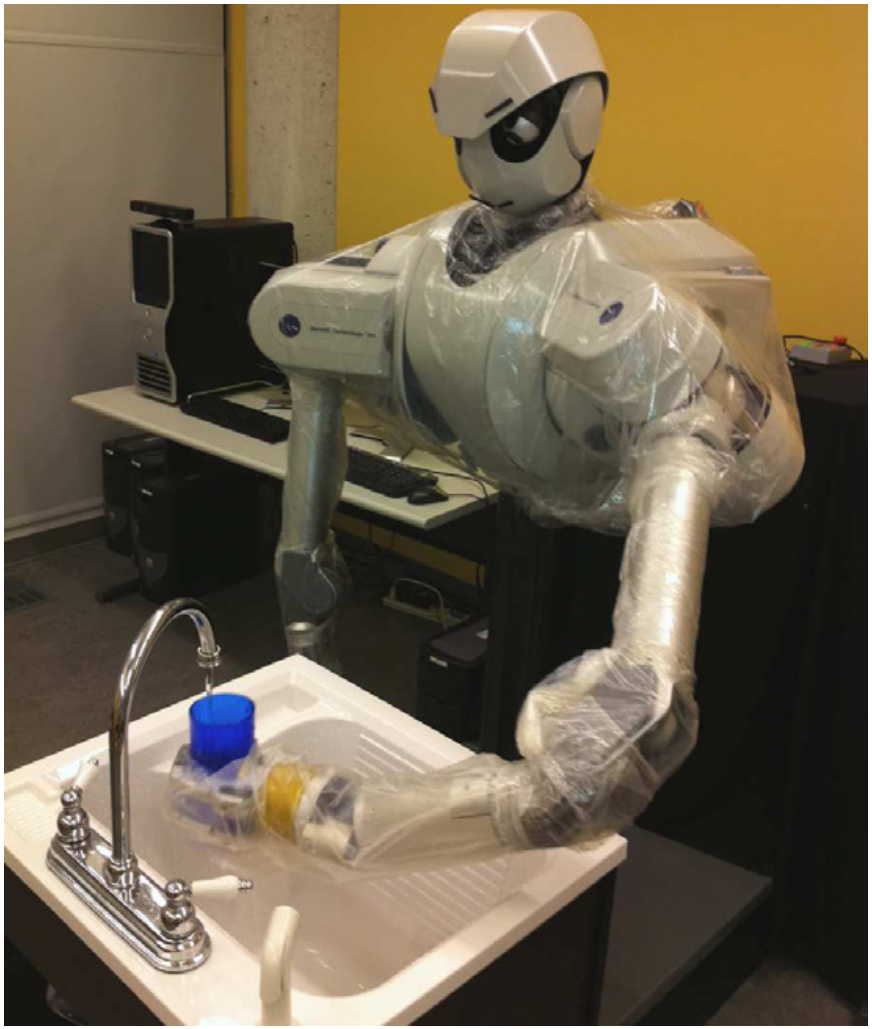

Water Use is Universal

Cooking

Cleaning

Gardening

How to Observe Water using Vision?

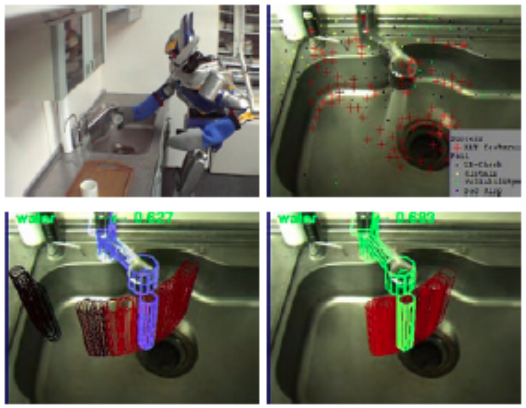

Okada *et al.*; 2009

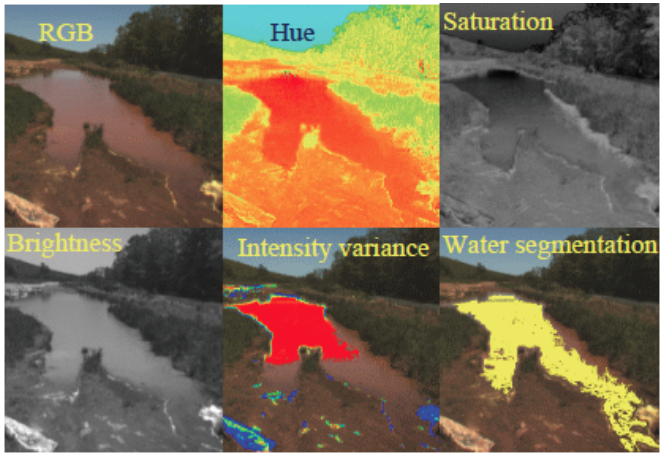

Rankin and Matthies; 2010

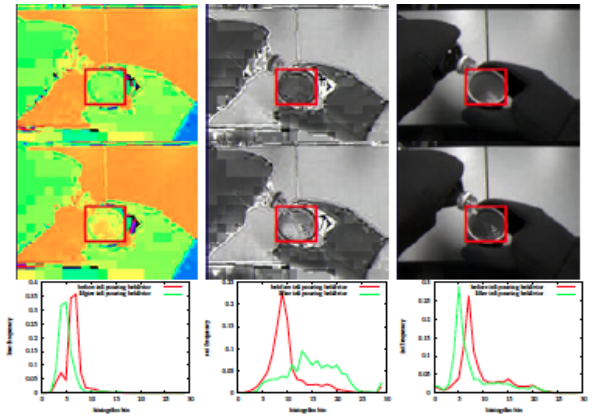

Okada *et al.*; 2009

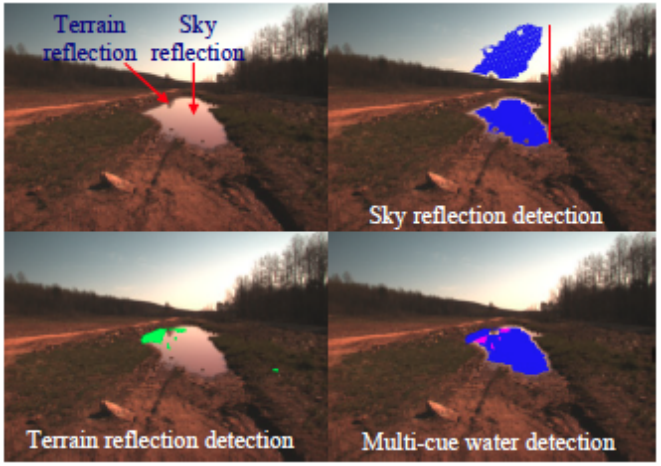

Rankin, Matthies, and Bellutta; 2011

Water is Hard to See

Water is Hard to See

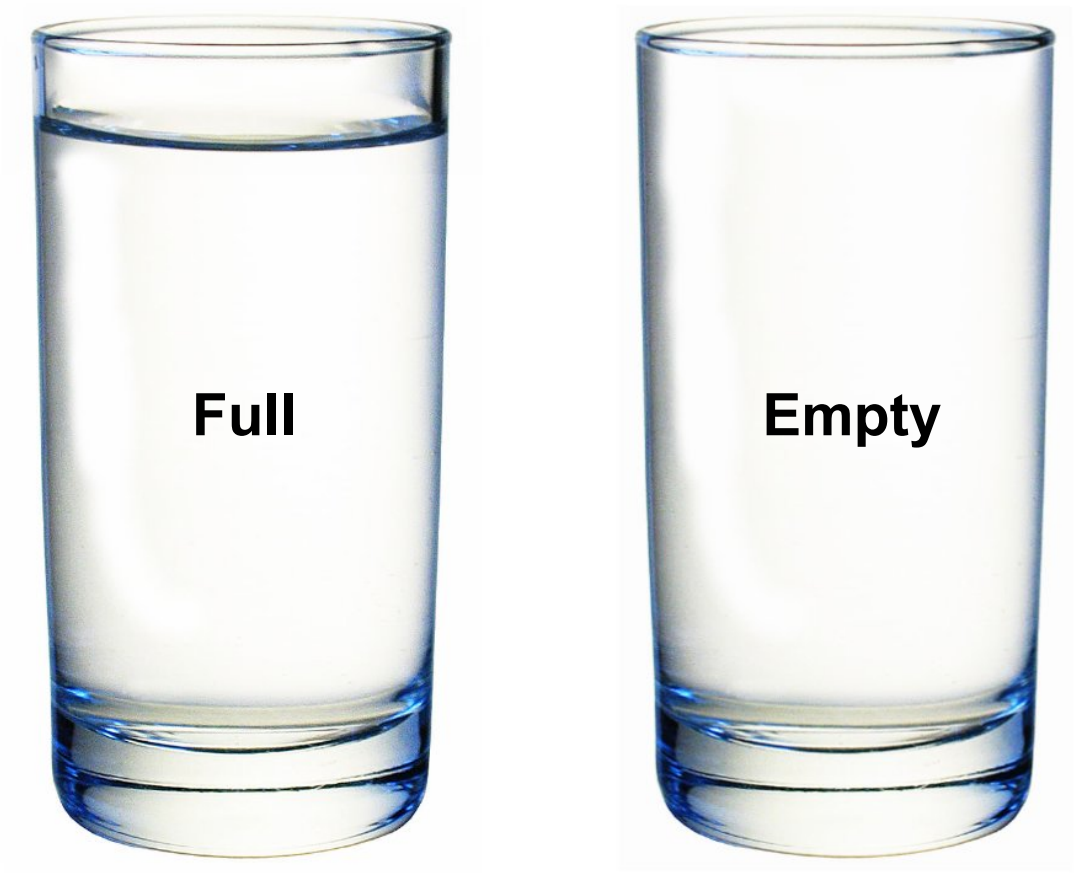

Some Objects That Hold Water Just Aren't Easy to Identify

Object Shape Can Be Deceptive






Research Question

What is a container?

Infant Interacting with a Container

Infant Interacting with a Non-Container

Previous Work

Vision and Audio

Griffith *et al.*; 2012

Audio and Proprioception

Sinapov *et al.*; 2011

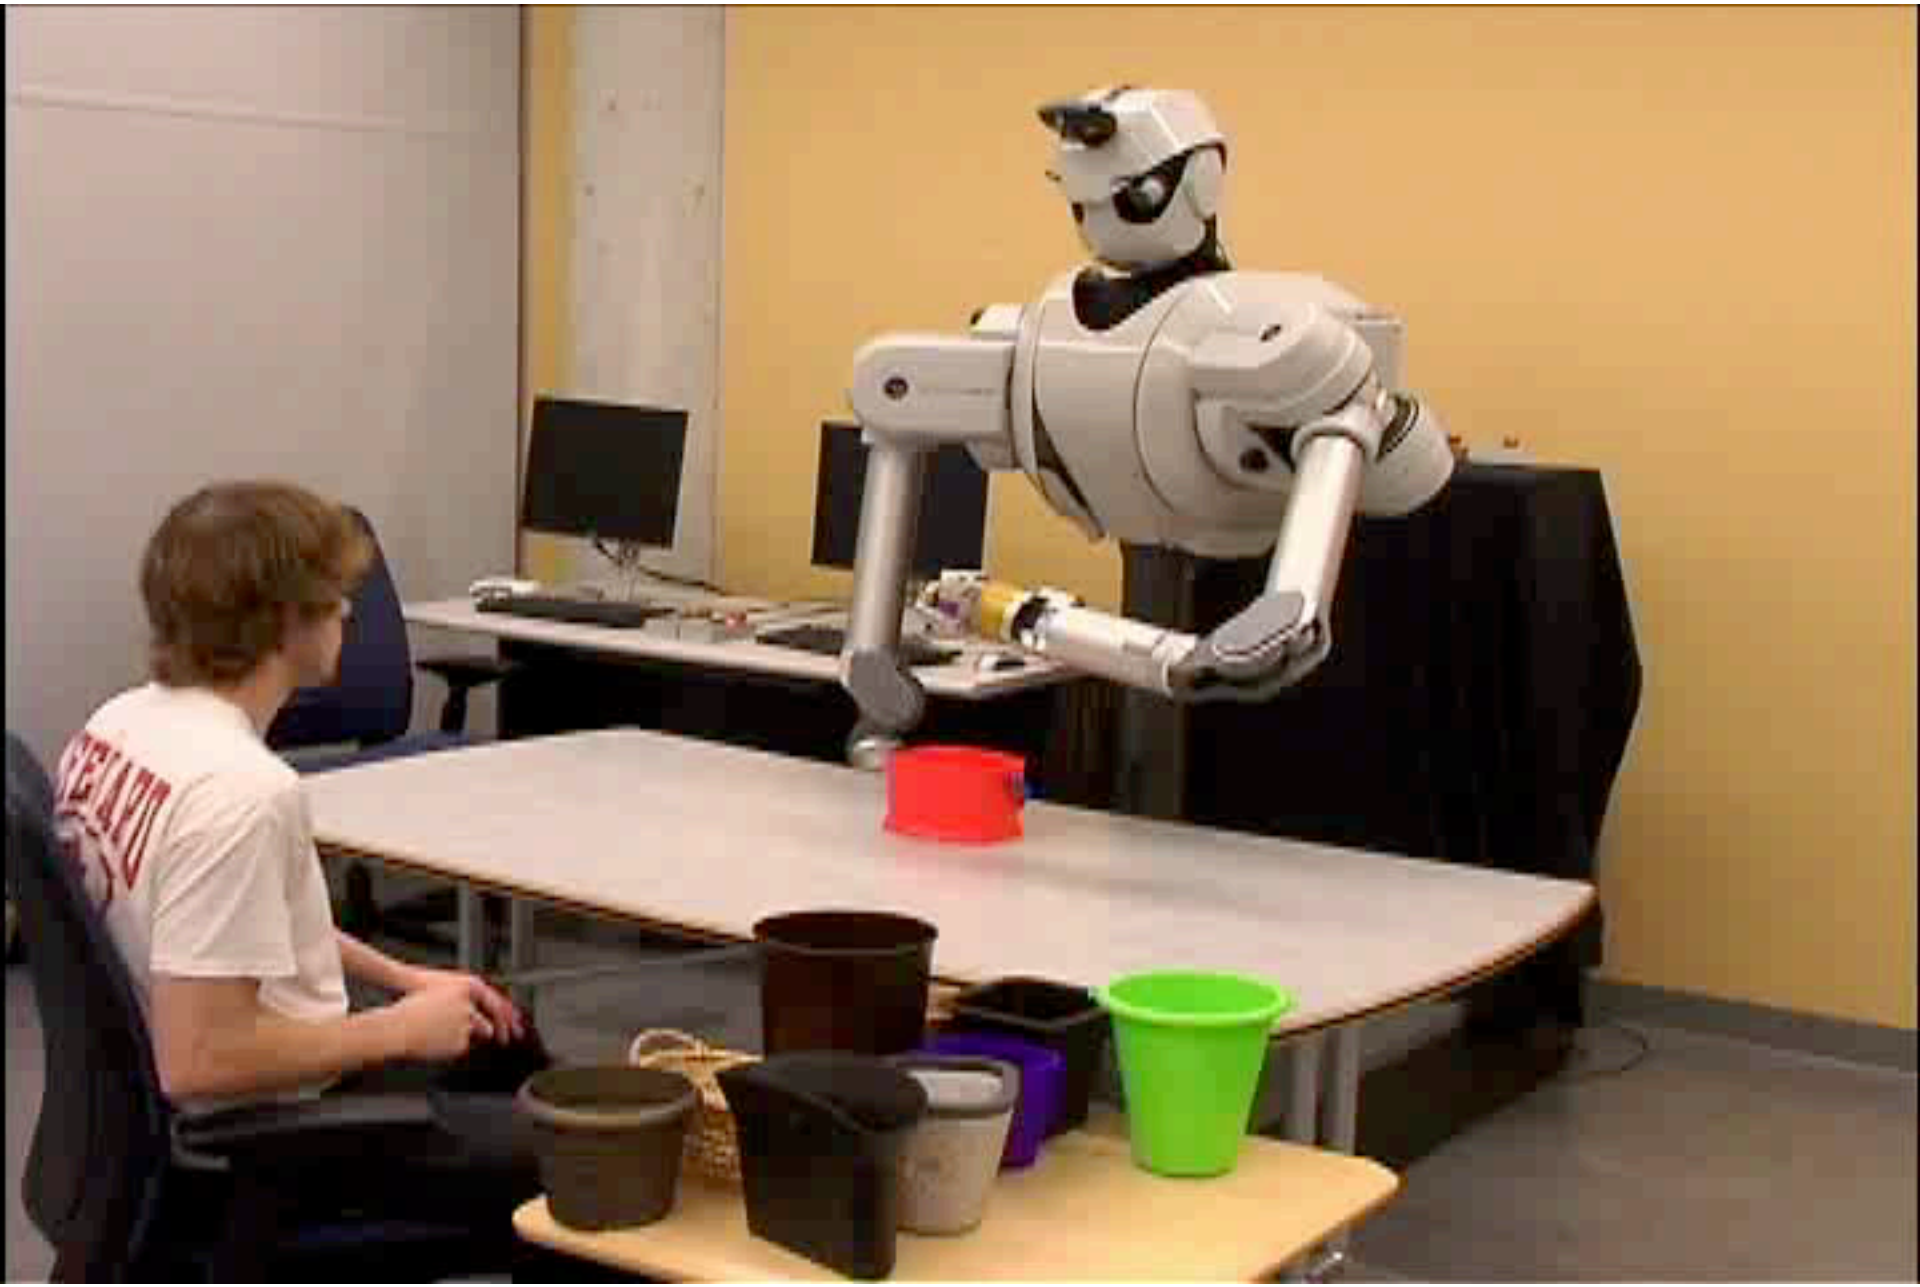

Learning Framework

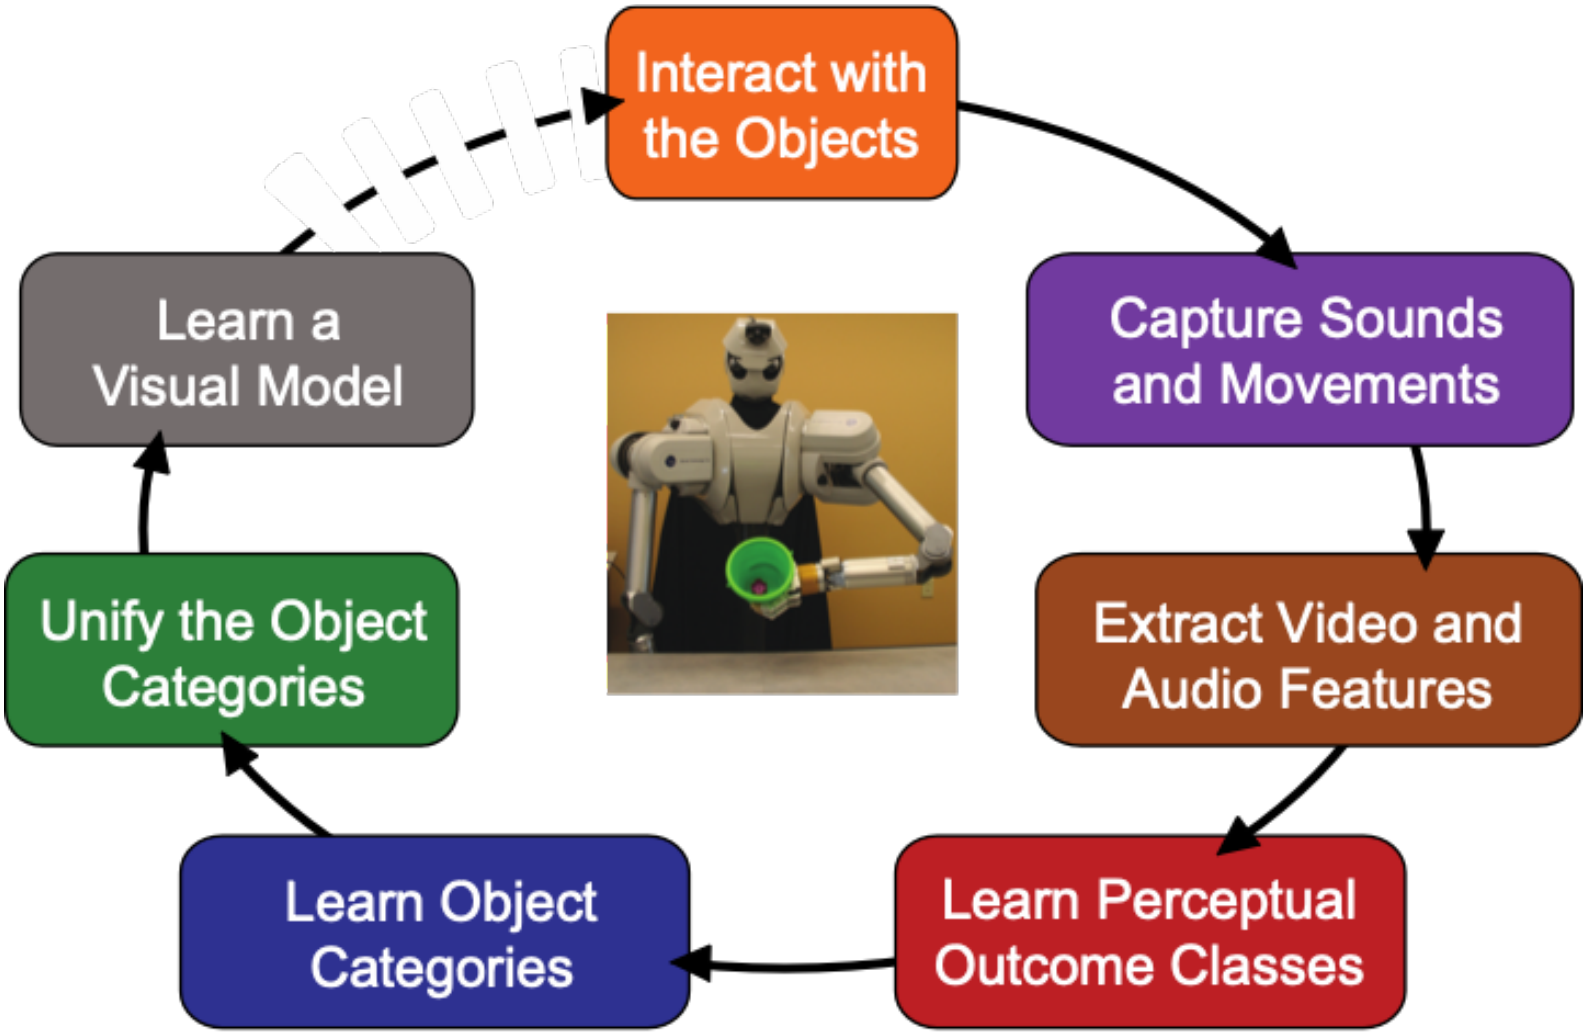

Learning Framework

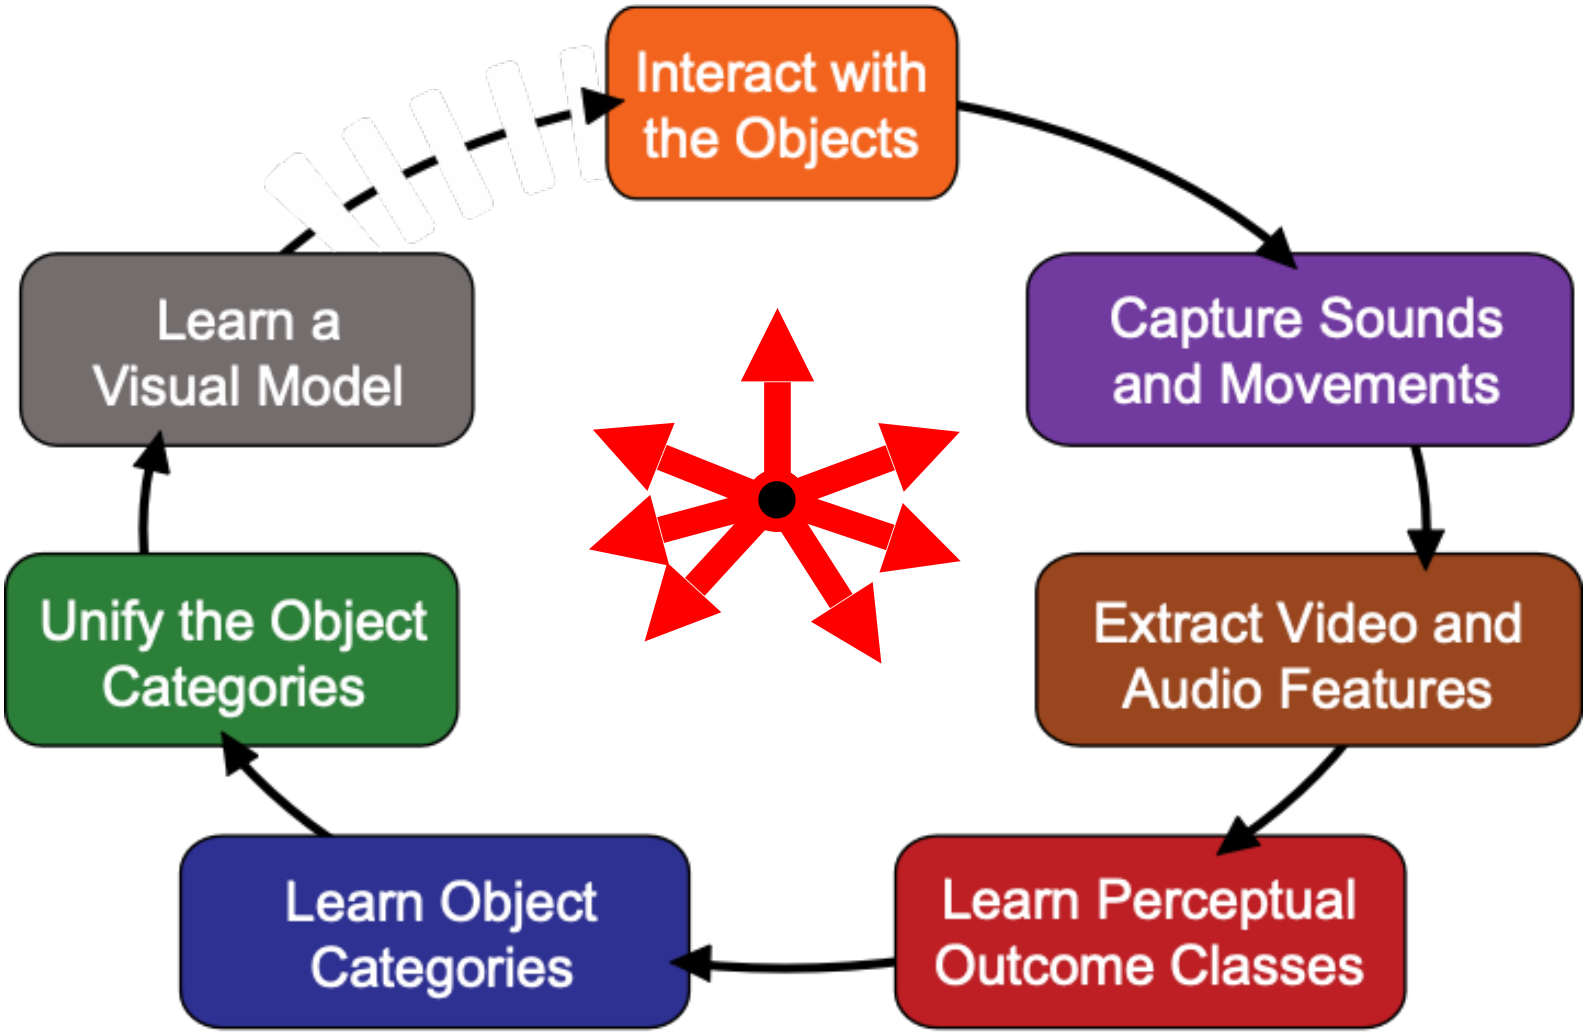

Unified Categorization

(derived from both sound and movement observations)

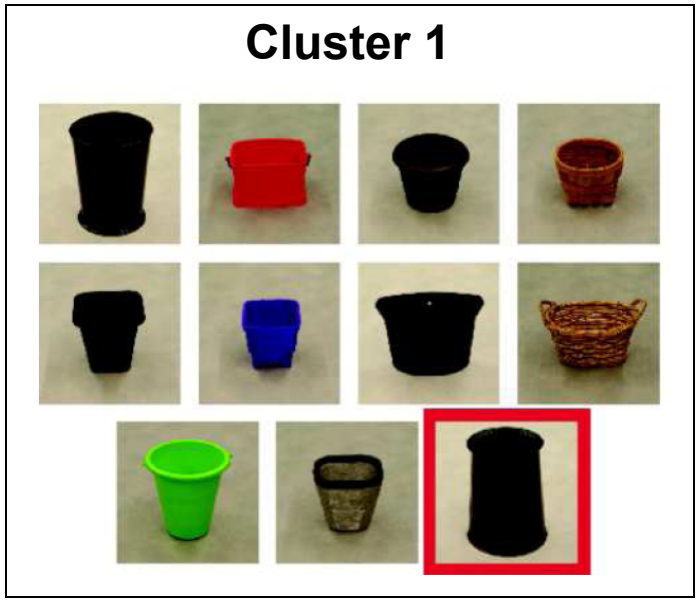

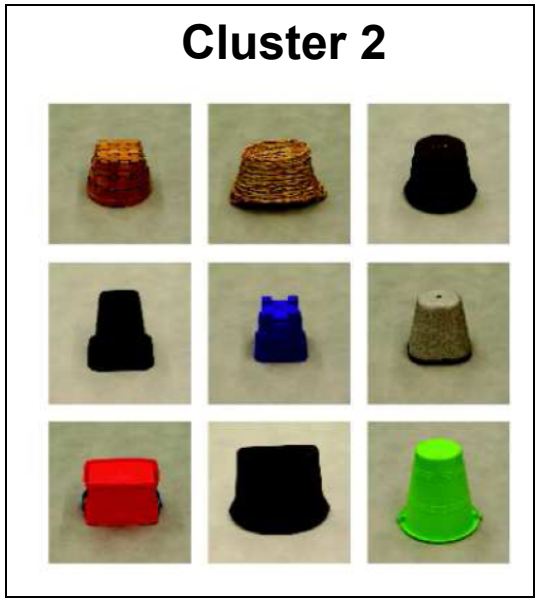

Extracted Visual Features

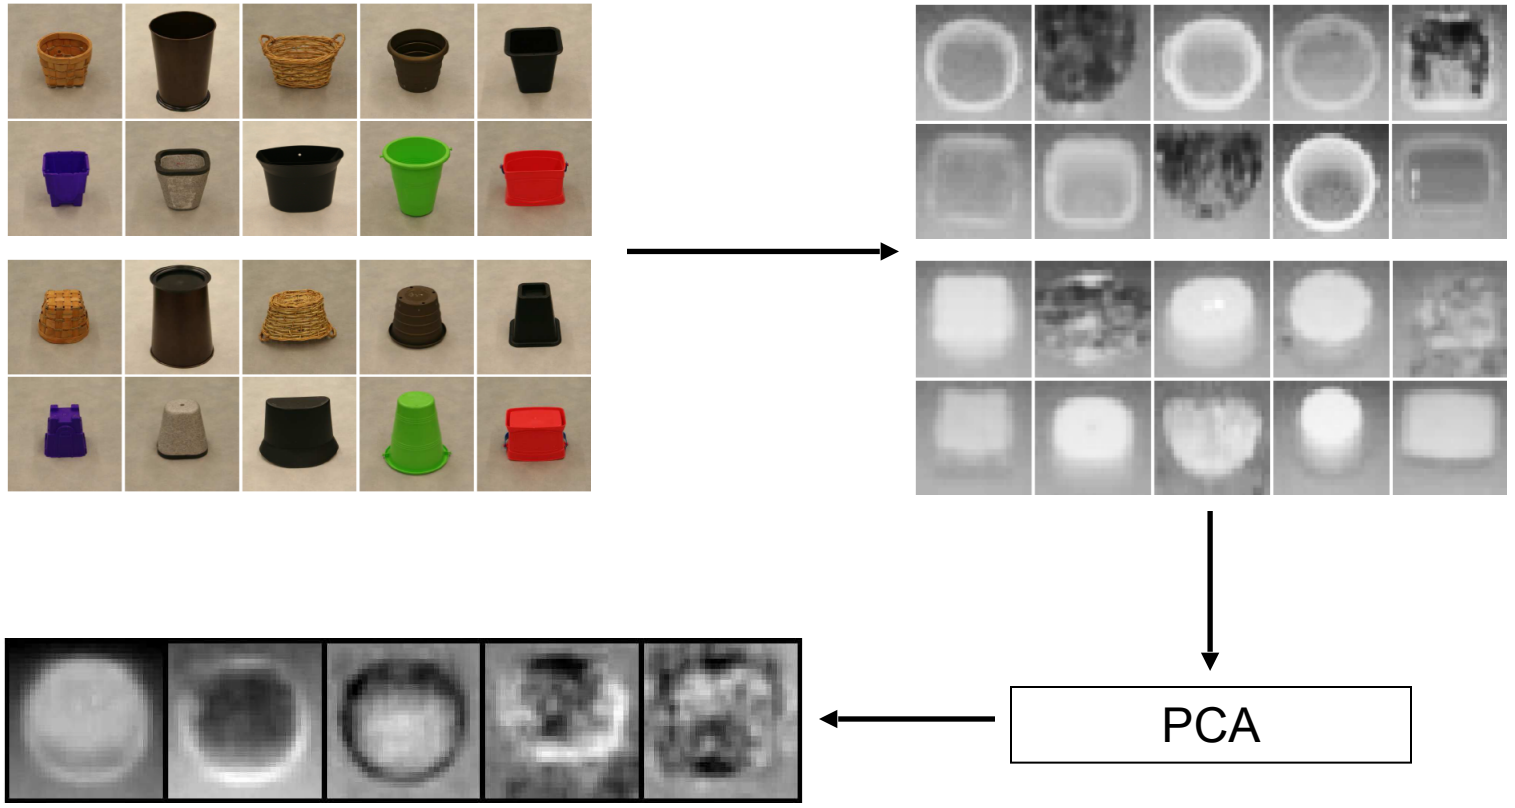

Example Visual Feature Activation

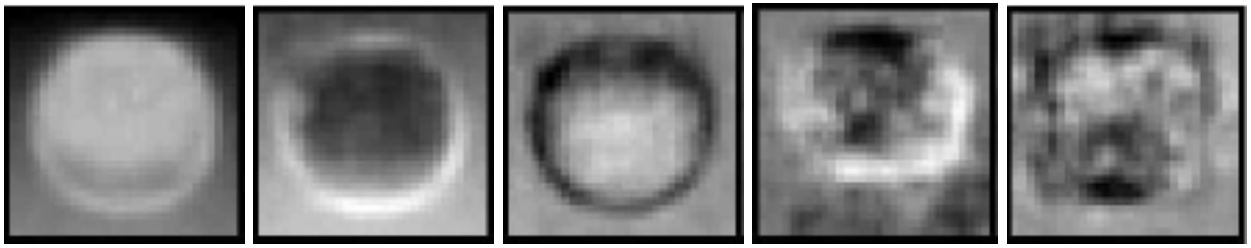

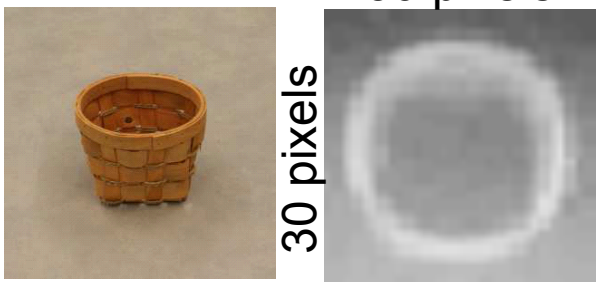

.19 .28 -.3 .02 .12

900 values → **PCA** → **5 values**

Example Container Feature Activation

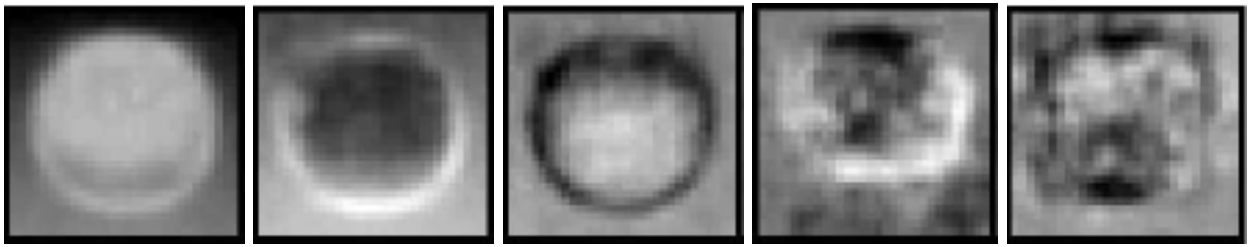

30 pixels

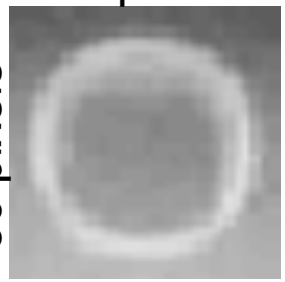

30 pixels

.19	.28	-.3	.02	.12
Strongly activated	Strongly activated	Strongly activated	Weakly activated	Weakly activated

The concave features are the most strongly activated

Example Non-Container Feature Activation

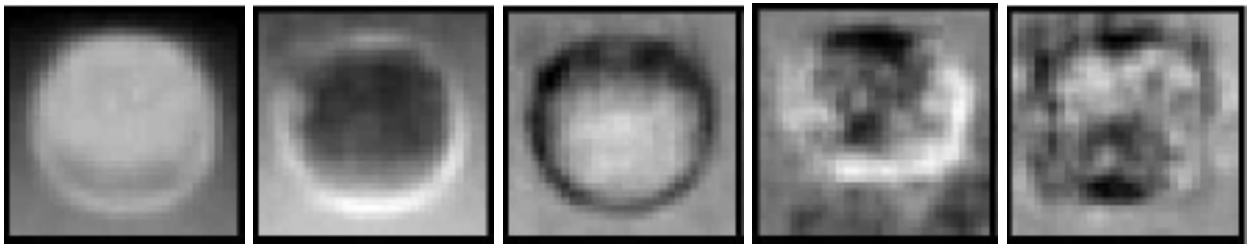

.31	-.14	.1	.04	-.08
Strongly activated	Weakly activated	Weakly activated	Weakly activated	Weakly activated

The convex feature is the most strongly activated

Classifying Novel Objects

Training Objects

					class label
0.19	0.28	-0.3	0.02	0.12	C_1
0.02	0.29	0.17	-0.03	0.67	C_1
0.22	0.2	-0.36	0.12	-0.06	C_1
0.2	0.32	0.0	0.04	0.0	C_1
-0.02	0.31	0.29	0.1	-0.22	C_1
0.18	0.31	0.0	-0.18	-0.13	C_1
0.28	0.13	-0.29	0.03	-0.03	C_1
0.02	0.27	0.41	-0.15	0.29	C_1
0.22	0.18	-0.4	0.0	0.07	C_1
0.04	0.35	0.09	0.04	-0.18	C_1

					class label
0.31	-0.14	0.1	0.04	-0.08	C_2
0.12	0.25	0.29	-0.29	-0.52	C_1
0.31	-0.14	0.09	0.04	0.07	C_2
0.3	-0.18	0.09	-0.09	-0.03	C_2
0.22	0.04	0.3	0.56	0.12	C_2
0.29	-0.14	0.11	-0.26	0.11	C_2
0.29	-0.18	0.14	0.0	0.08	C_2
0.25	-0.15	0.16	0.44	-0.14	C_2
0.24	-0.17	0.04	-0.48	0.09	C_2
0.29	-0.14	-0.02	-0.08	-0.02	C_2

Novel Object


0.0 0.07 -0.2 -0.3 0.1

kNN

Cluster 1

Classifying Novel Objects

Training Objects

					class label
0.19	0.28	-0.3	0.02	0.12	C_1
0.02	0.29	0.17	-0.03	0.67	C_1
0.22	0.2	-0.36	0.12	-0.06	C_1
0.2	0.32	0.0	0.04	0.0	C_1
-0.02	0.31	0.29	0.1	-0.22	C_1
0.18	0.31	0.0	-0.18	-0.13	C_1
0.28	0.13	-0.29	0.03	-0.03	C_1
0.02	0.27	0.41	-0.15	0.29	C_1
0.22	0.18	-0.4	0.0	0.07	C_1
0.04	0.35	0.09	0.04	-0.18	C_1
0.31	-0.14	0.1	0.04	-0.08	C_2
0.12	0.25	0.29	-0.29	-0.52	C_1
0.31	-0.14	0.09	0.04	0.07	C_2
0.3	-0.18	0.09	-0.09	-0.03	C_2
0.22	0.04	0.3	0.56	0.12	C_2
0.29	-0.14	0.11	-0.26	0.11	C_2
0.29	-0.18	0.14	0.0	0.08	C_2
0.25	-0.15	0.16	0.44	-0.14	C_2
0.24	-0.17	0.04	-0.48	0.09	C_2
0.29	-0.14	-0.02	-0.08	-0.02	C_2

Novel Object


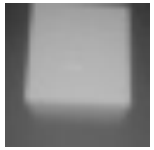

0.2 -0.2 0.0 0.05 -0.2

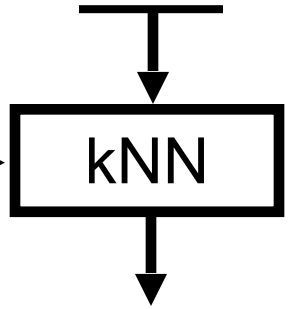

Cluster 2

Classification Results

Novel Containers

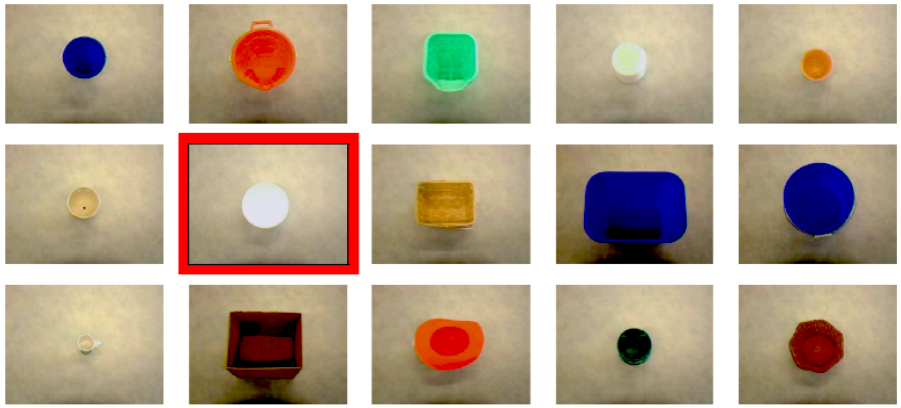

Novel Non-Containers

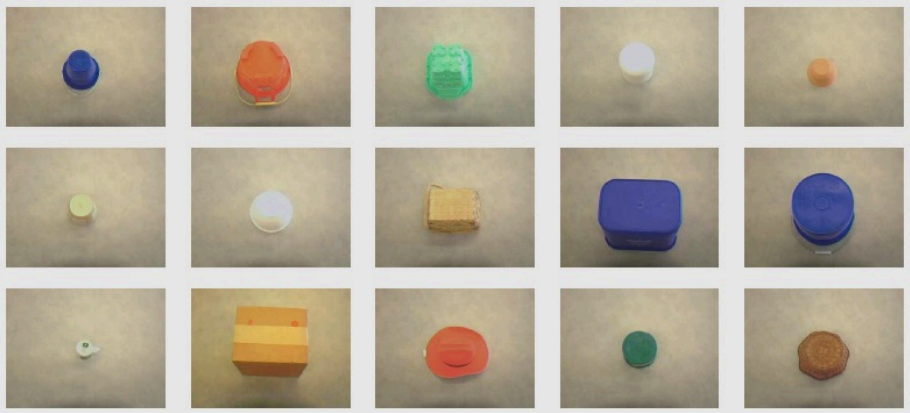

Infant Playing in the Sink

Video of the Experiments

Interactive Object Categorization

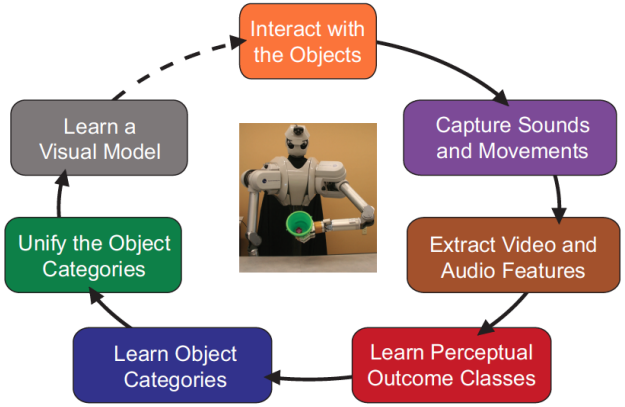

Griffith, Sinapov, Sukhoy, and Stoytchev; 2012

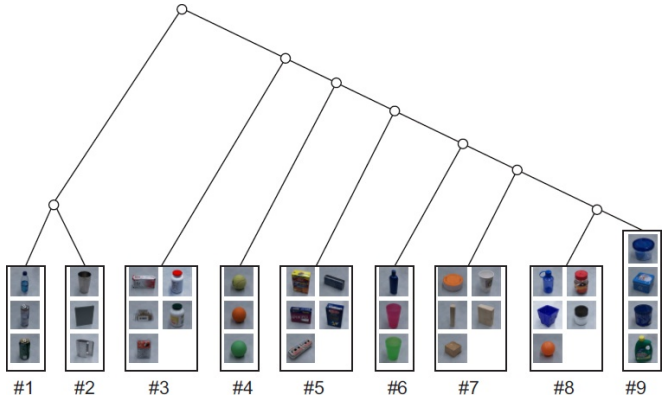

Sinapov and Stoytchev; 2009

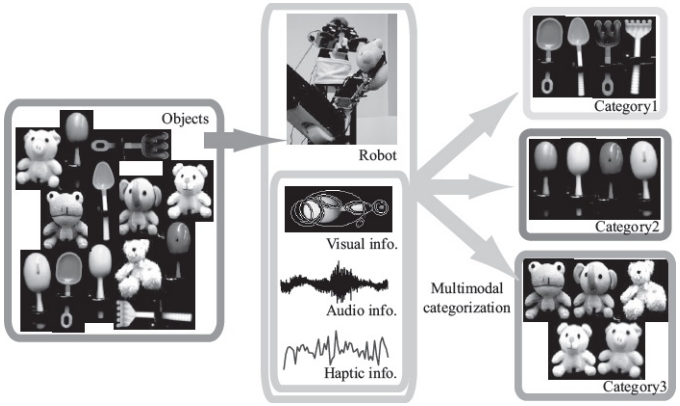

Nakamura, Nagai, and Iwahashi; 2007

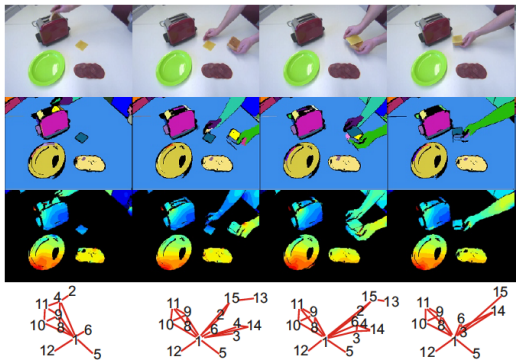

Aksoy *et al.*; 2010

Slosh-Free Control of Containers

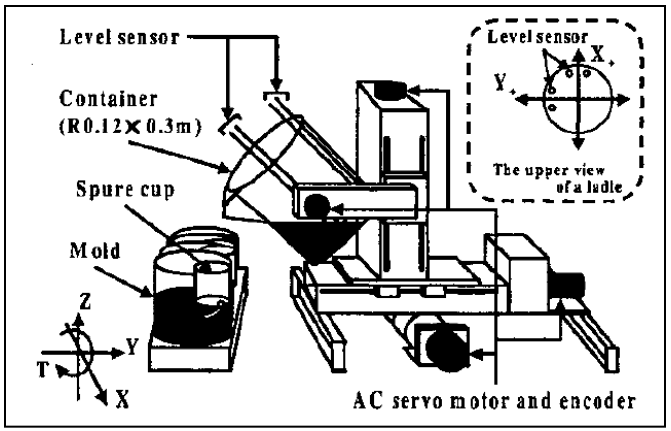

Yano *et al.*; 2001

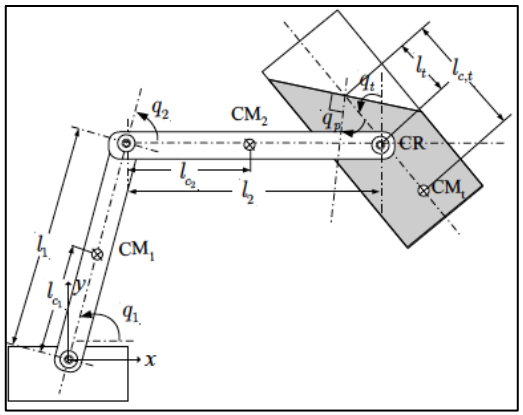

Tzamtzi, Koumboulis, Kouvakas; 1997

Feddema *et al.*; 1997

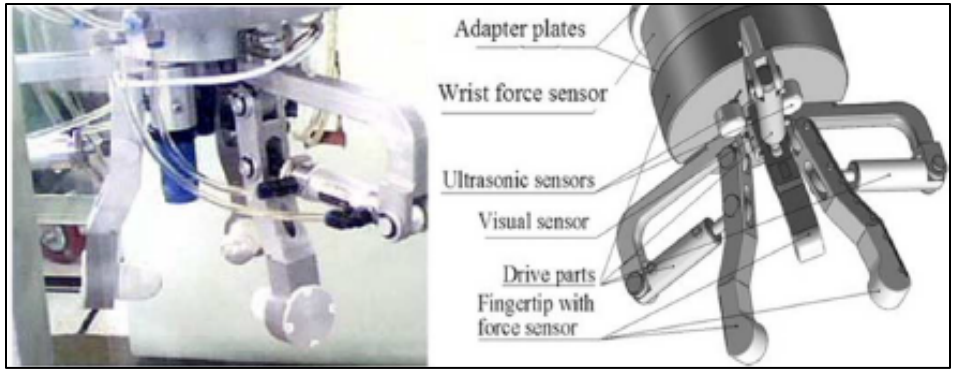

Liang, Zhang, Song, and Ge; 2010

Pouring Liquid into a Container

Cakmak and Thomaz; 2011

Okada *et al.*; 2009

Kim *et al.*; 2009

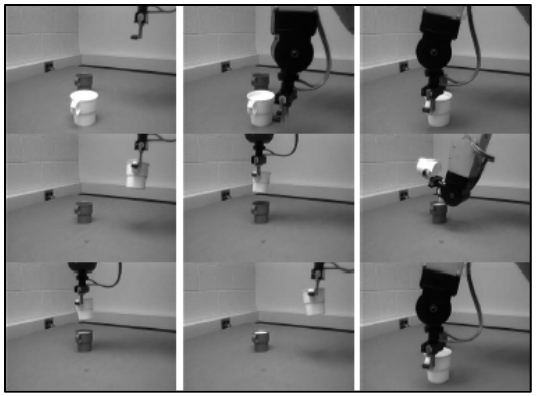

Hwang and Weng; 1997

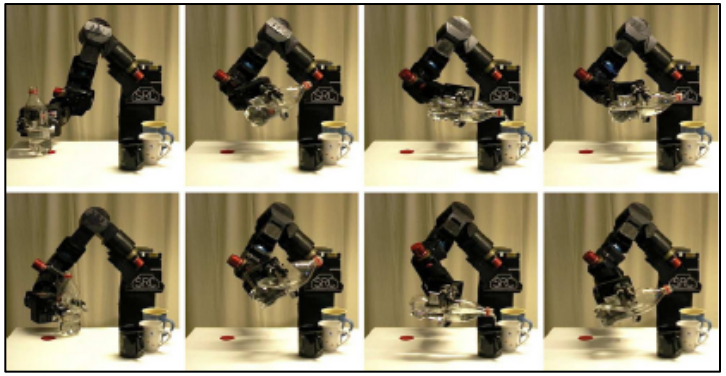

Pastor *et al.*; 2009

Sink

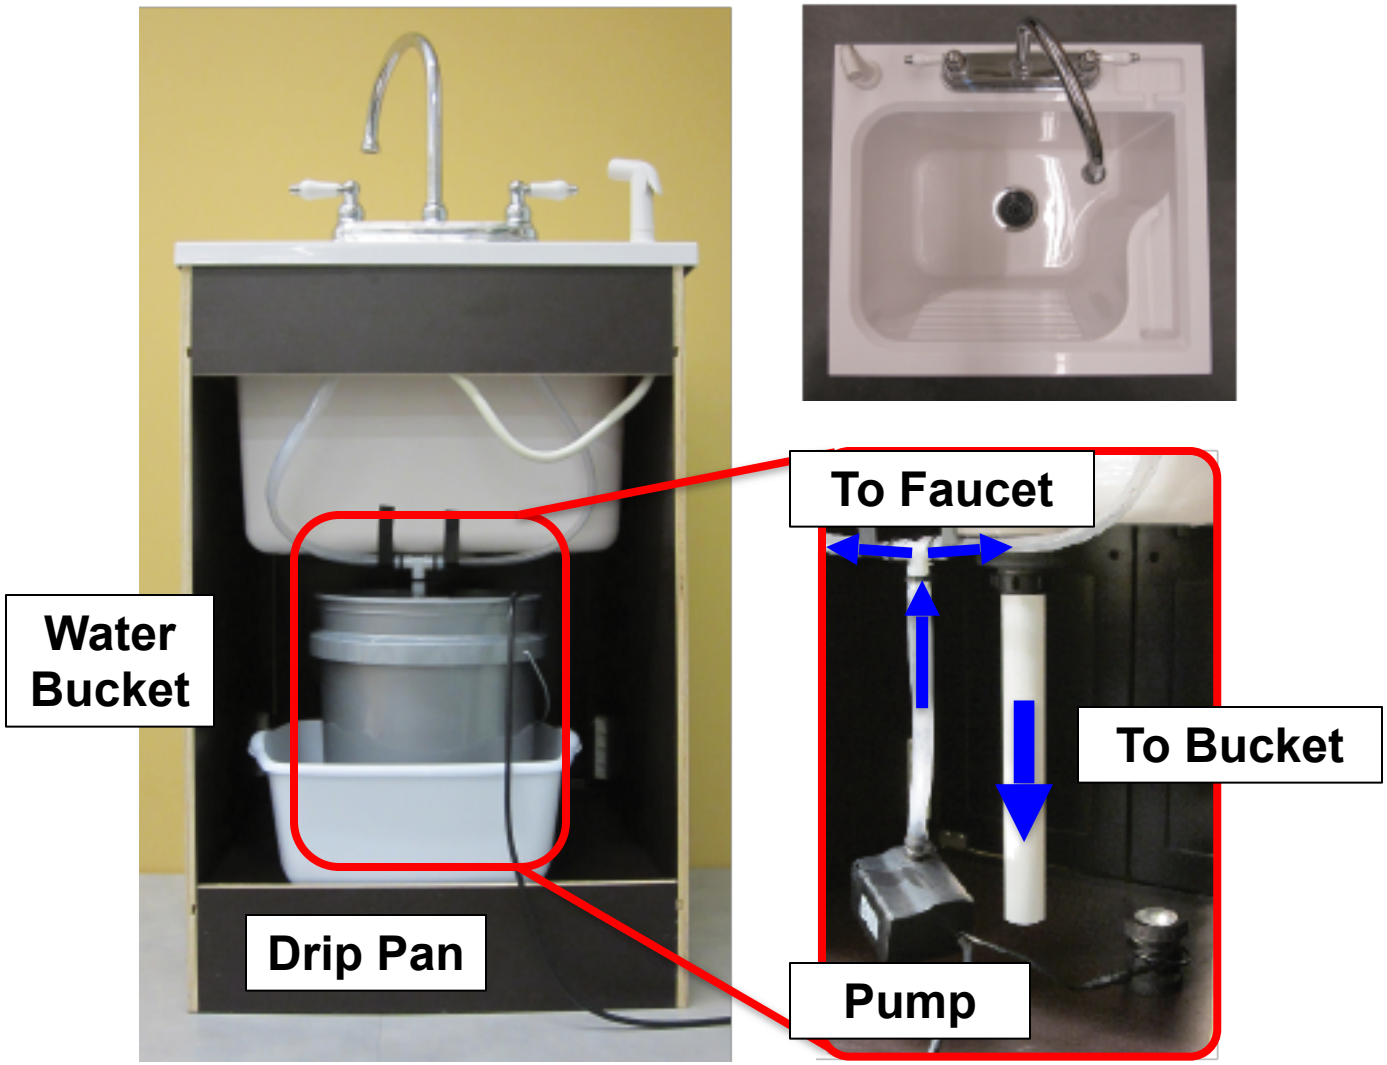

Can a Robot Categorize These Objects Using Audio and Proprioception?

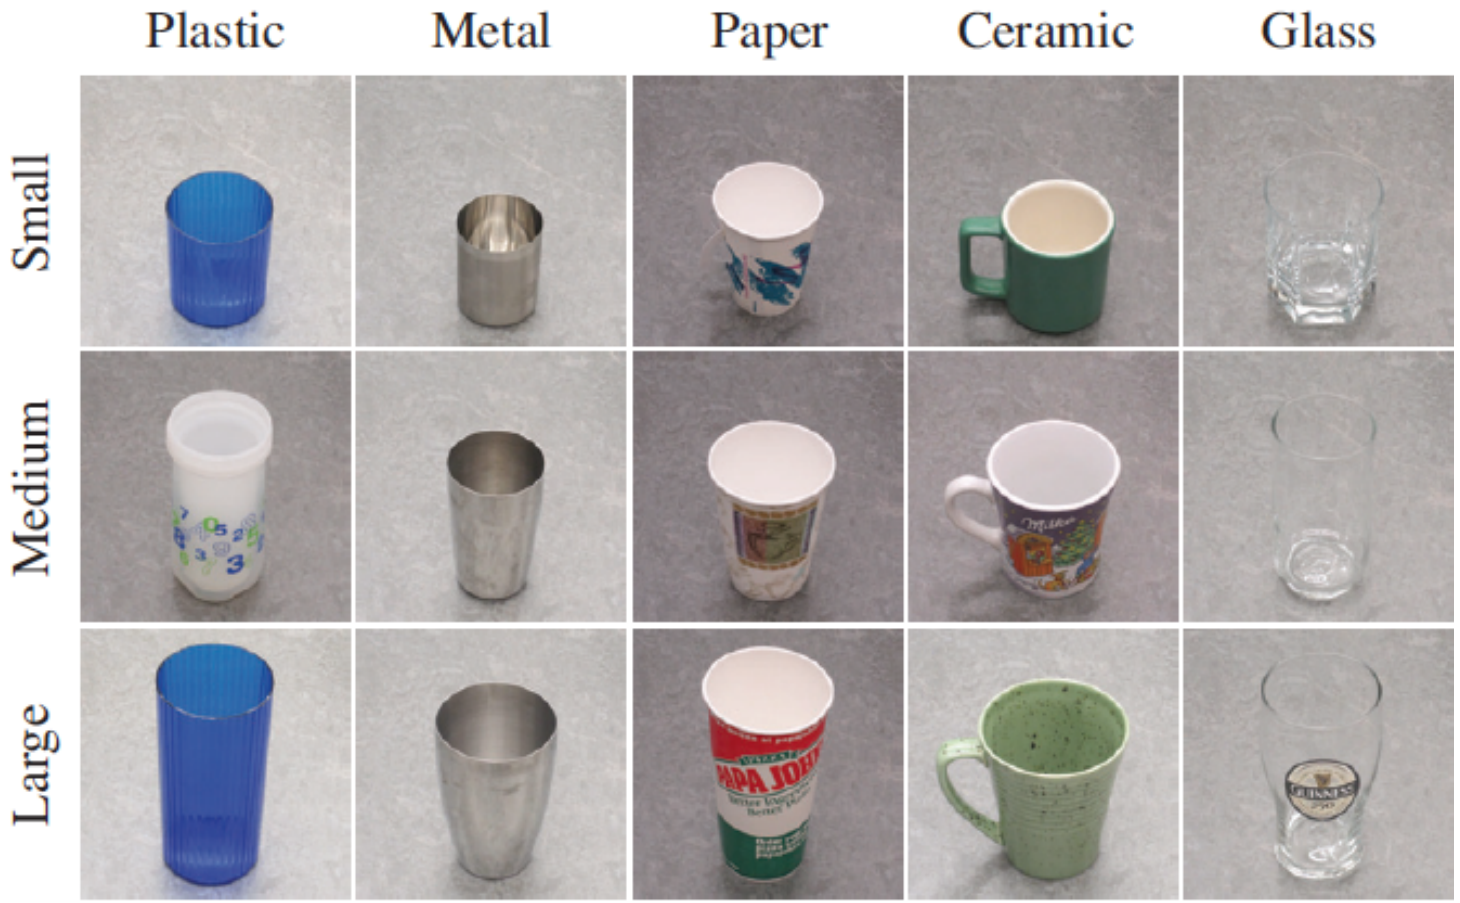

Learning Framework

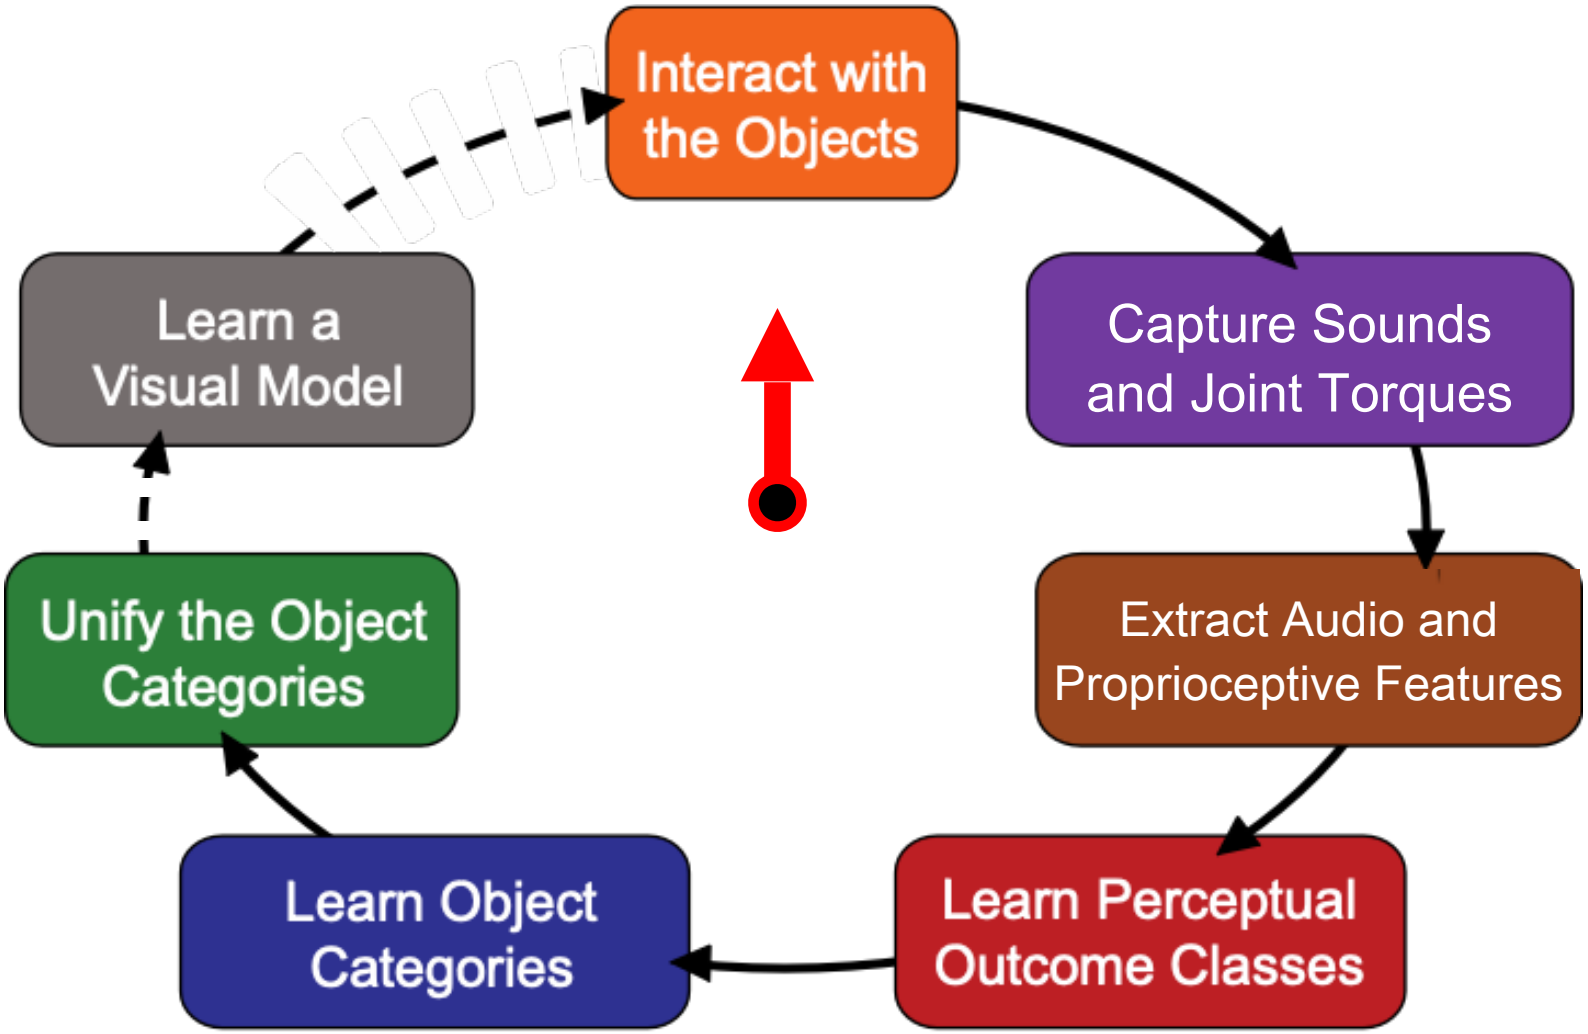

Behaviors

Don’t Fry This At Home

Waterguard Cast and Skin Protector

Learning Framework

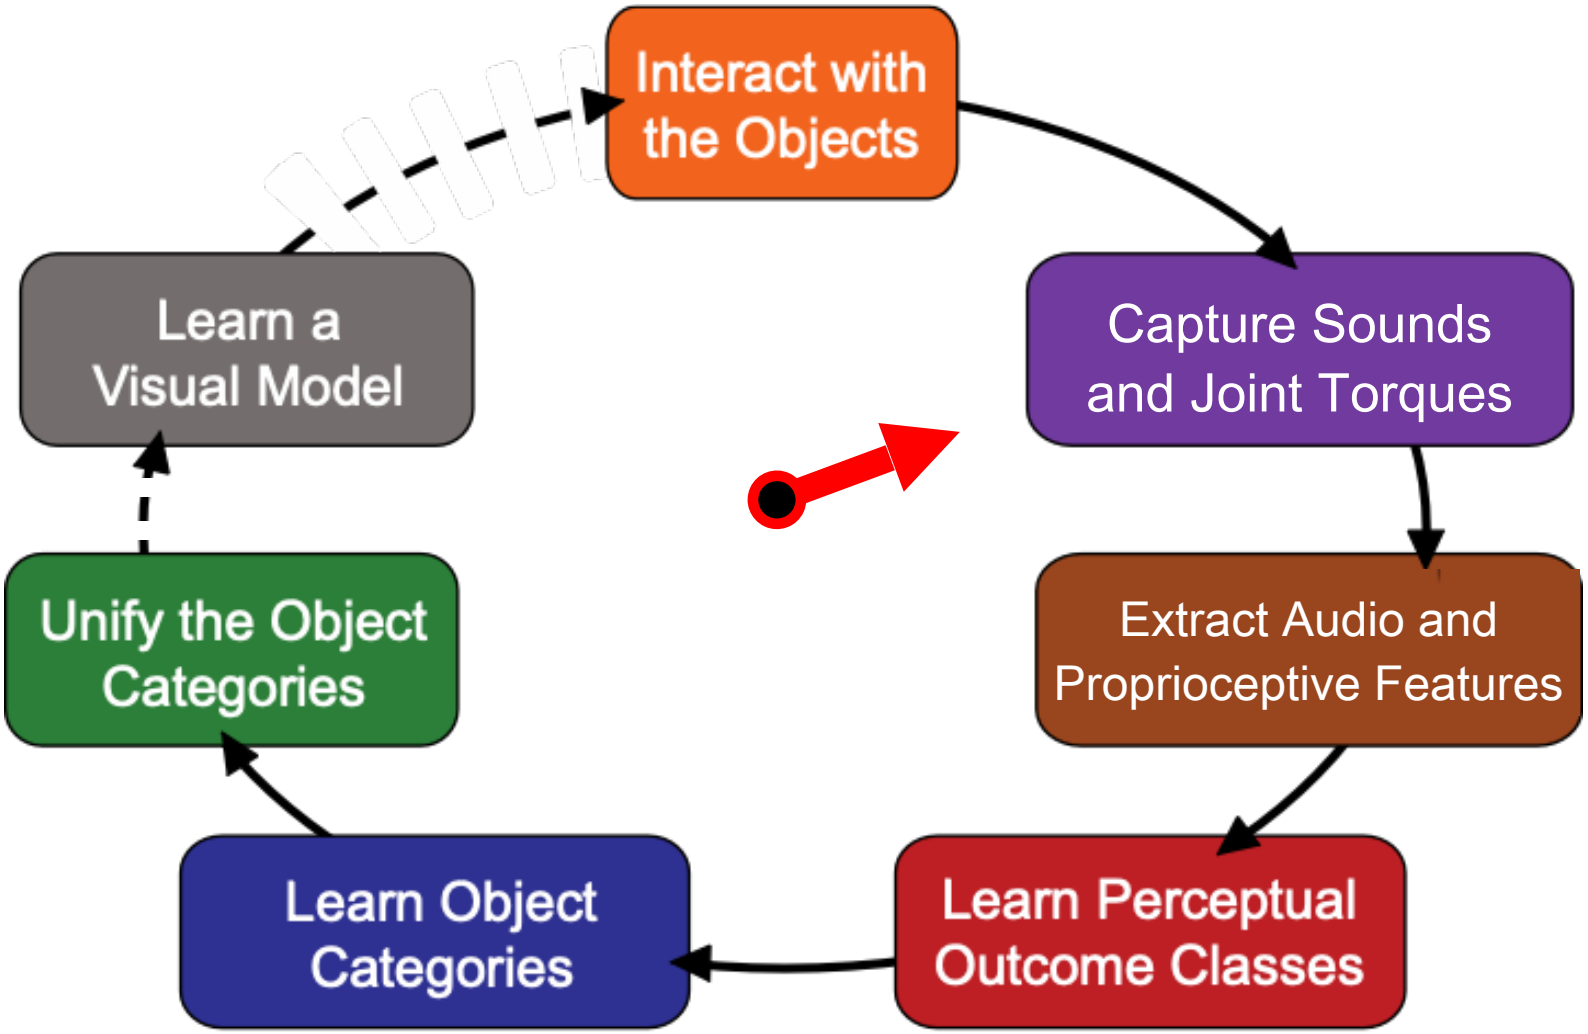

Data Collection

5 behaviors x 10 trials x 15 objects x 2 object poses

1,500 behavioral interactions

(6 hours of interaction)

Learning Framework

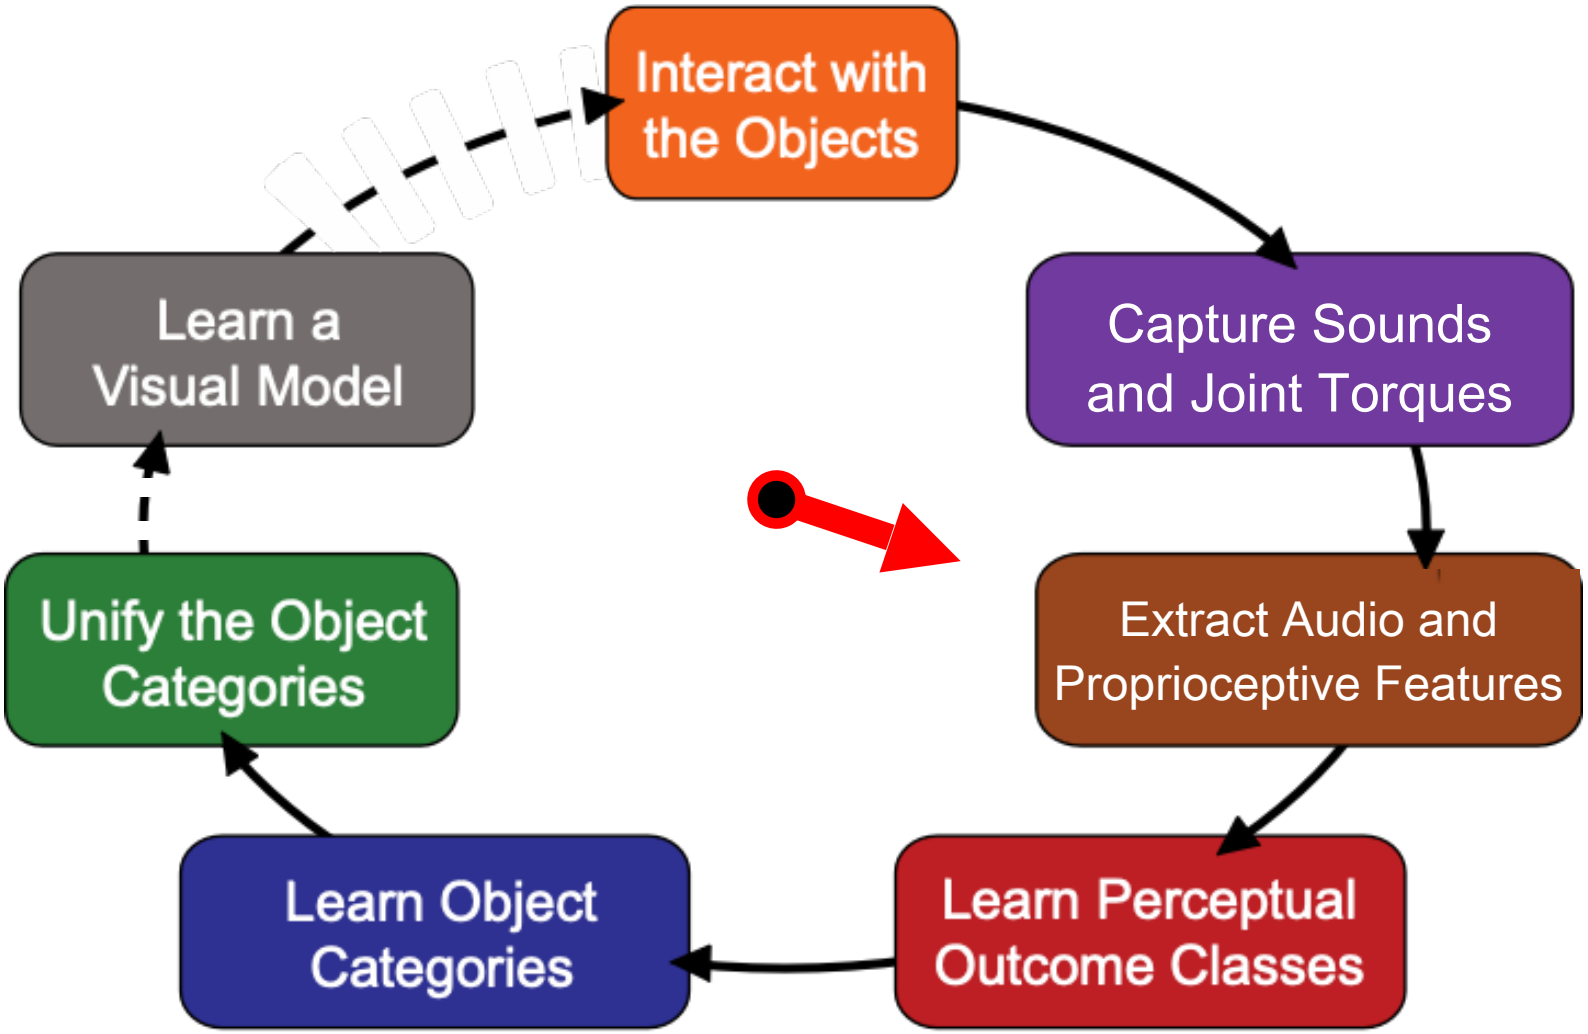

Audio Preprocessing

Behavior Execution:
(up and down)

WAV file recorded:

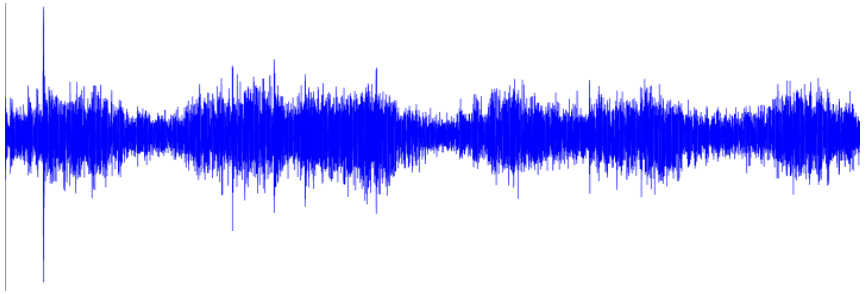

Discrete Fourier Transform:

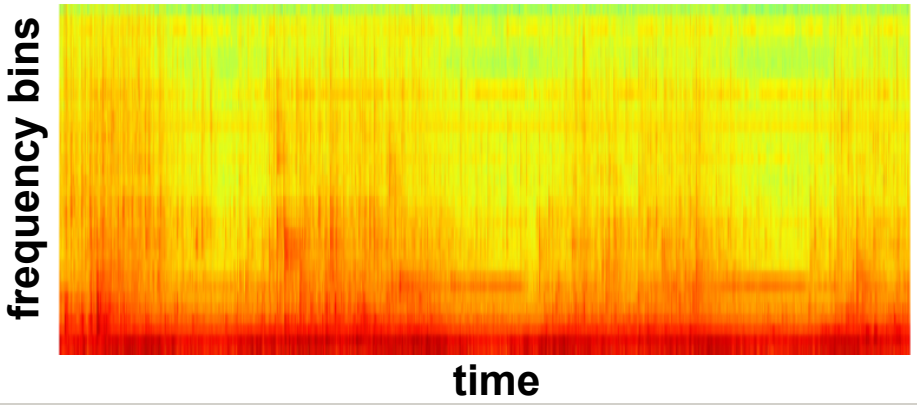

Unsupervised Feature Extraction

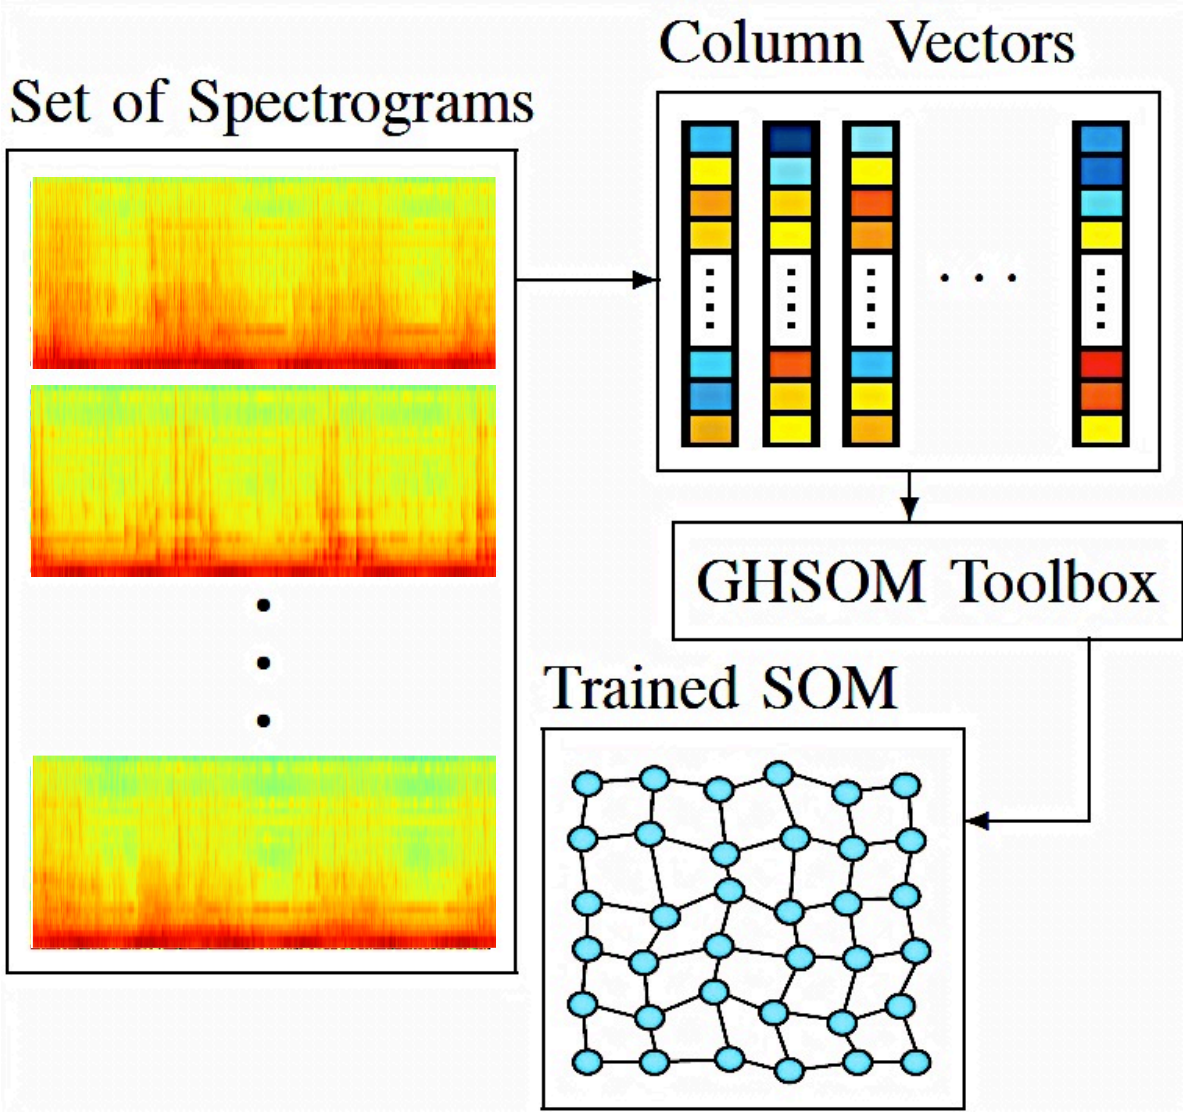

Convert the Spectrogram to a State Sequence Using a SOM

Spectrogram:

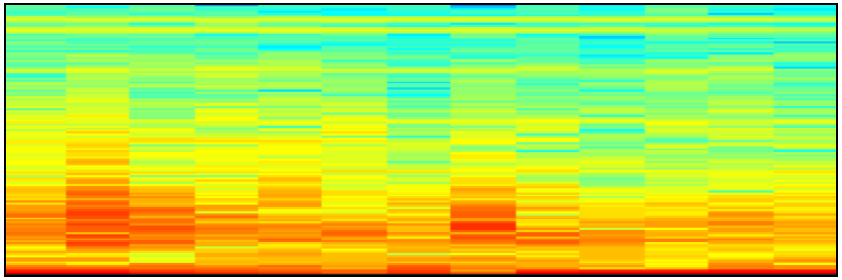

Self Organizing Map:

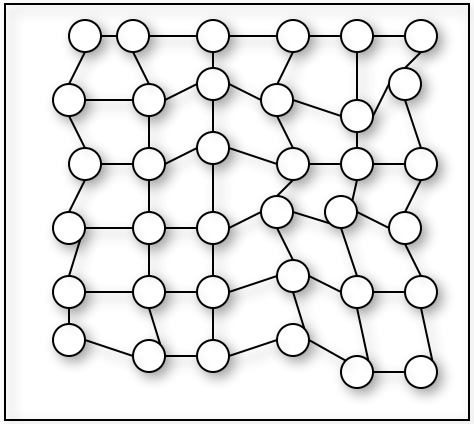

Convert the Spectrogram to a State Sequence Using a SOM

Spectrogram:

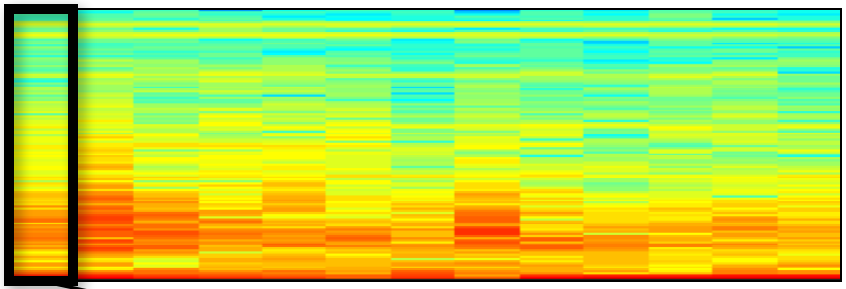

Self Organizing Map:

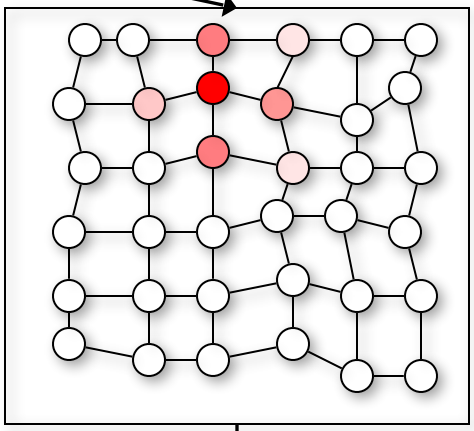

State Sequence:

A_i: (3,5) ->

Convert the Spectrogram to a State Sequence Using a SOM

Spectrogram:

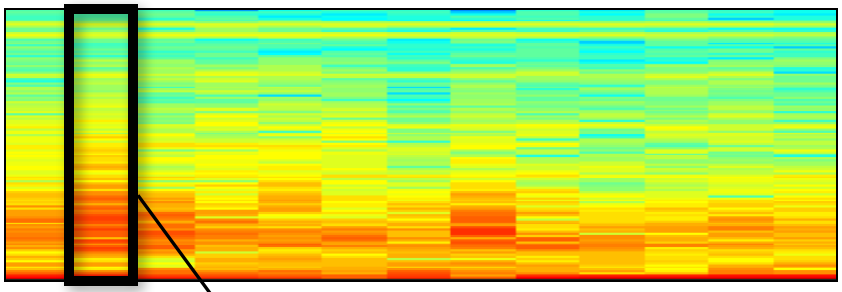

Self Organizing Map:

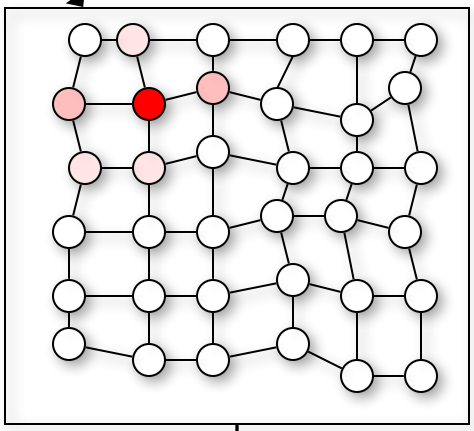

State Sequence:

A_i: (3,5) -> (2,5) ->

Convert the Spectrogram to a State Sequence Using a SOM

Spectrogram:

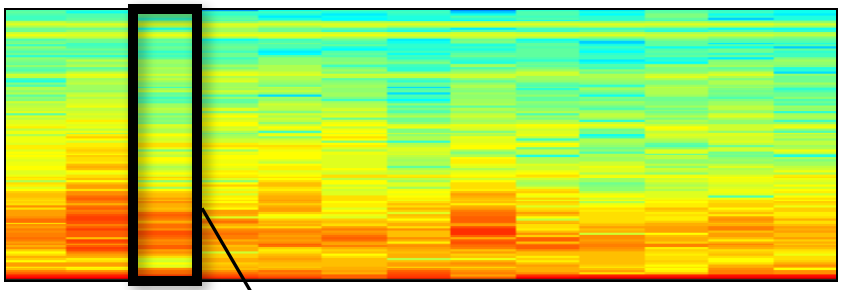

Self Organizing Map:

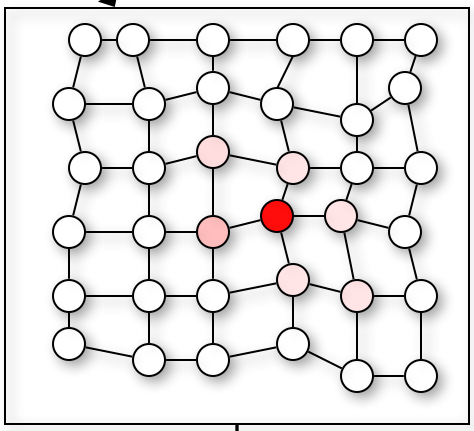

State Sequence:

A_i: (3,5) -> (2,5) -> (4,3) -> ….

Acoustic Feature Extraction

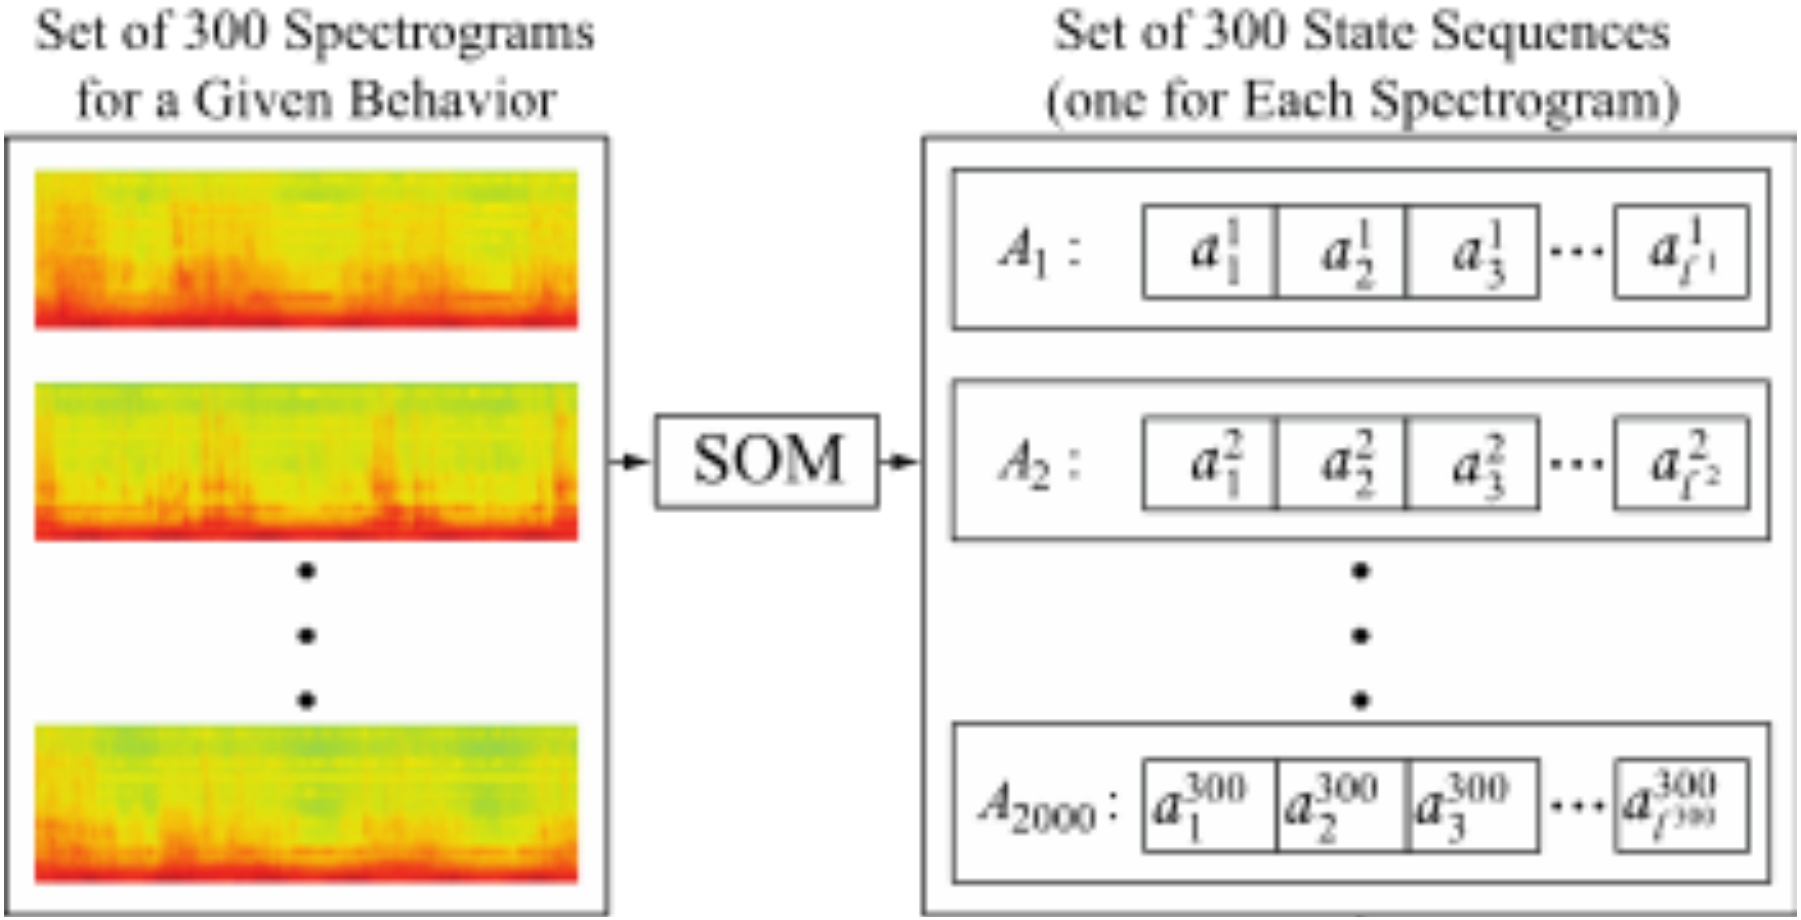

Proprioceptive Feature Extraction

Behavior Execution:
(in and out)

Joint Torque Sequence:

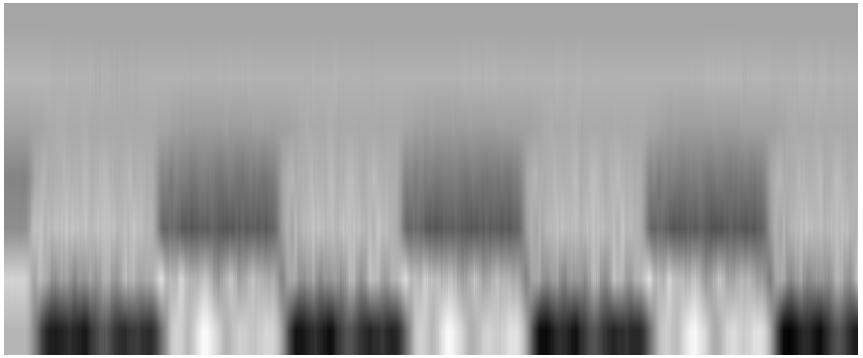

Unsupervised Feature Extraction

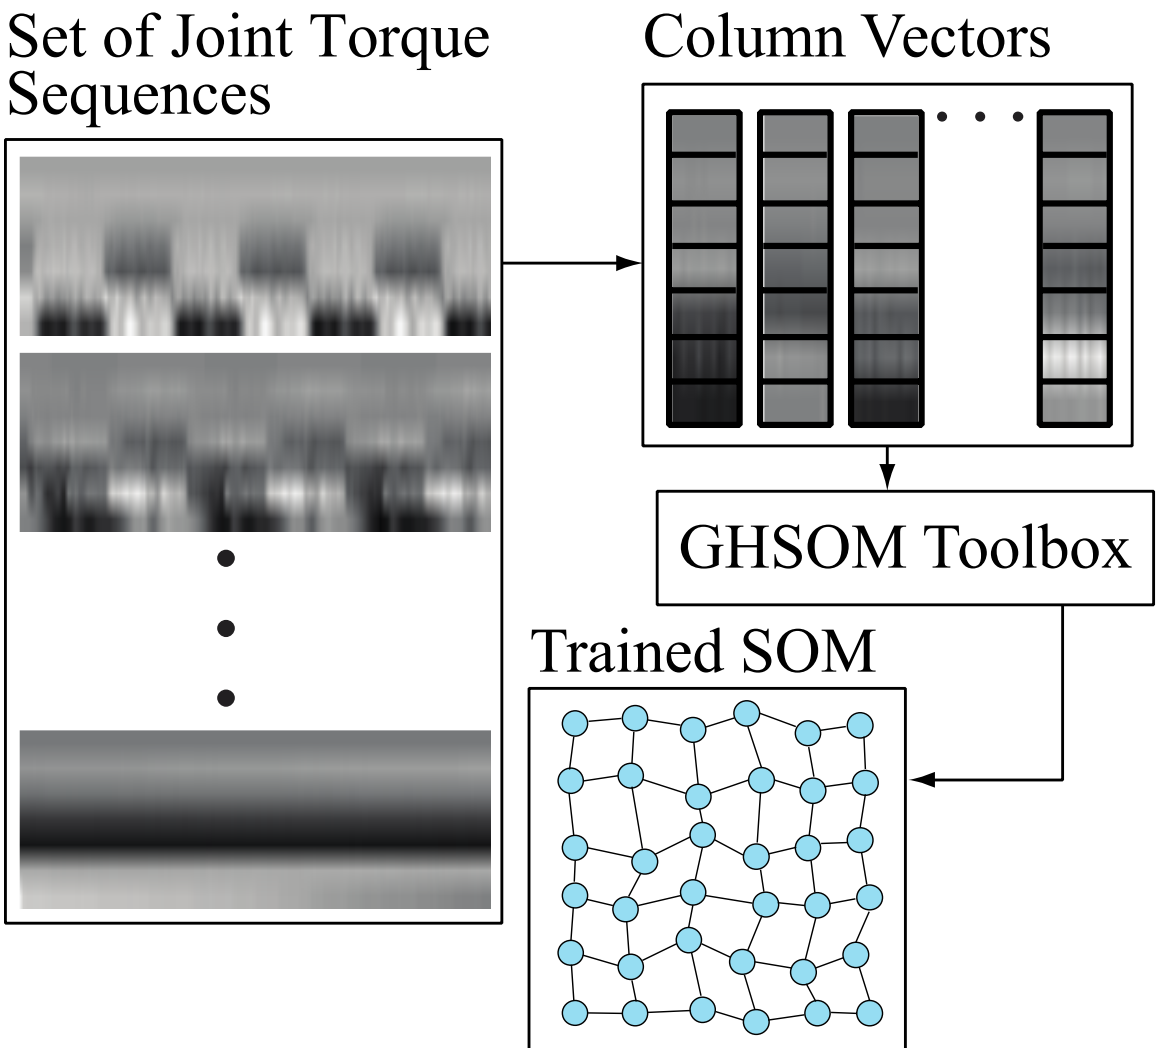

Convert the Joint Torque Sequence to a State Sequence Using a SOM

Joint Torque Sequence:

Self Organizing Map:

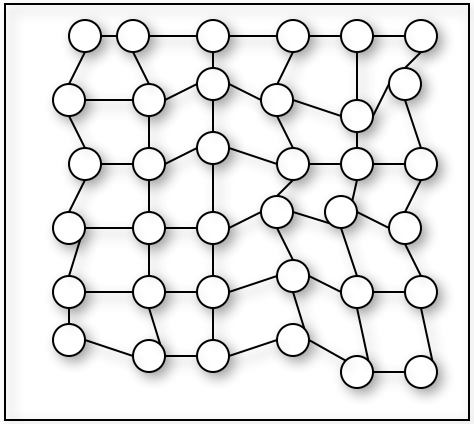

Convert the Joint Torque Sequence to a State Sequence Using a SOM

Joint Torque Sequence:

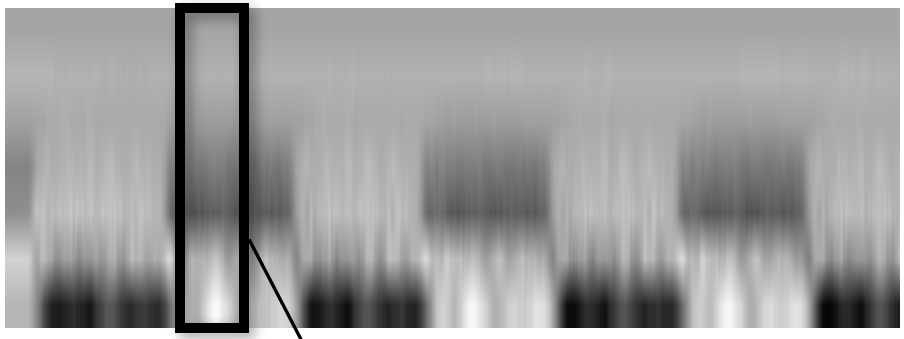

Self Organizing Map:

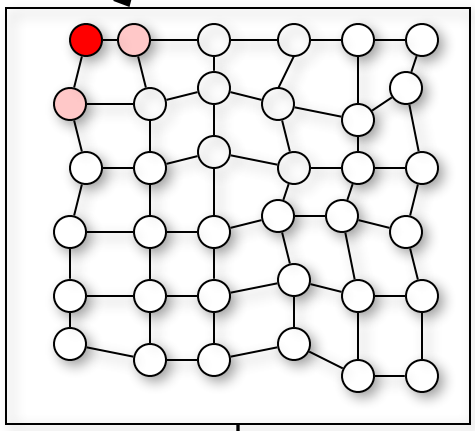

State Sequence:

P_i: (1,6) ->

Convert the Joint Torque Sequence to a State Sequence Using a SOM

Joint Torque Sequence:

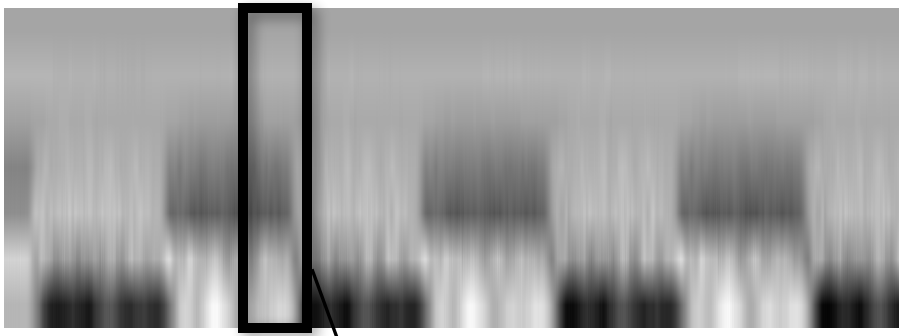

Self Organizing Map:

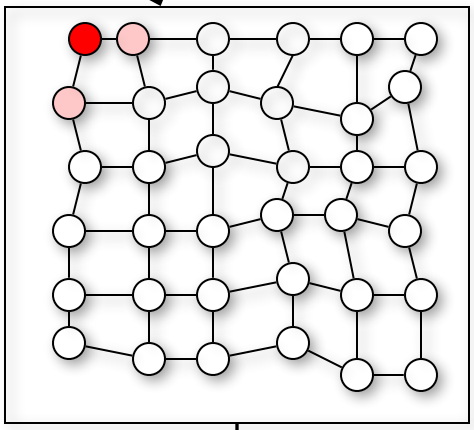

State Sequence:

P_i: (1,6) -> (1,6) ->

Convert the Joint Torque Sequence to a State Sequence Using a SOM

Joint Torque Sequence:

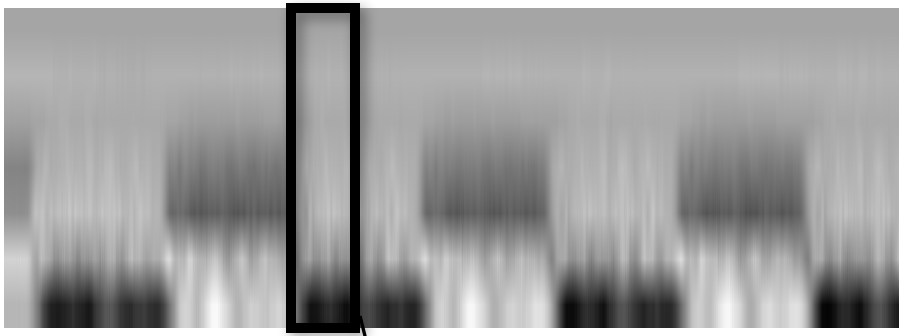

Self Organizing Map:

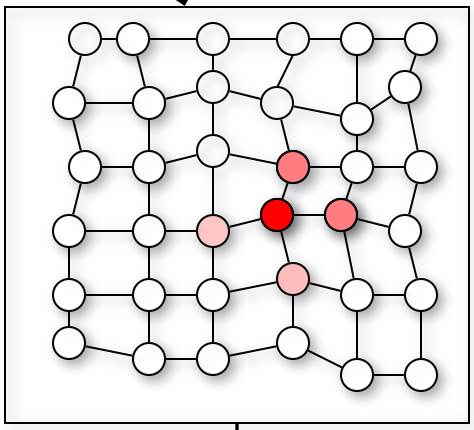

State Sequence:

P_i: (1,6) -> (1,6) -> (4,3) -> ….

Proprioception Feature Extraction

Set of 300 Joint Torque Sequences for a Given Behavior

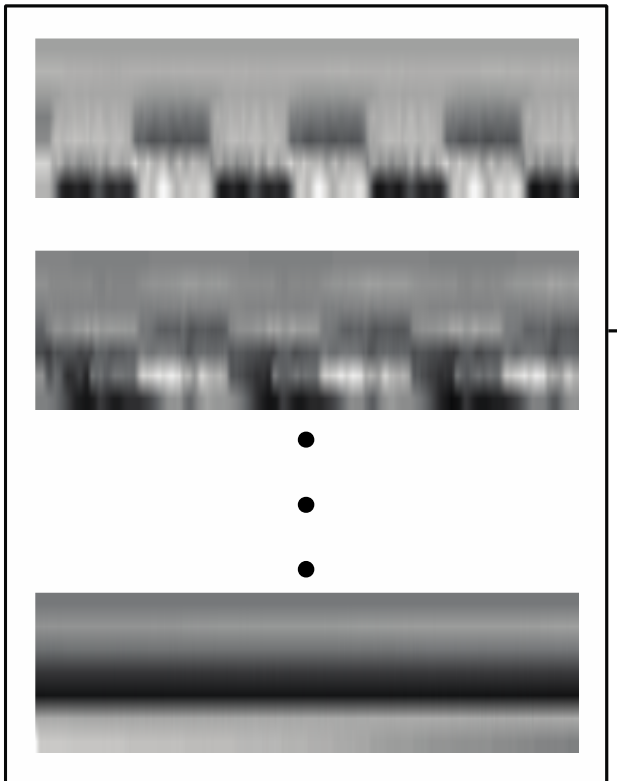

SOM

Set of 300 State Sequences (one for Each Joint Torque Sequence)

P_1 : p_1^1 p_2^1 p_3^1 ⋯ $p_{l_1}^1$

P_2 : p_1^2 p_2^2 p_3^2 ⋯ $p_{l_2}^2$

P_{300} : p_1^{300} p_2^{300} p_3^{300} ⋯ $p_{l_{300}}^{300}$

Learning Framework

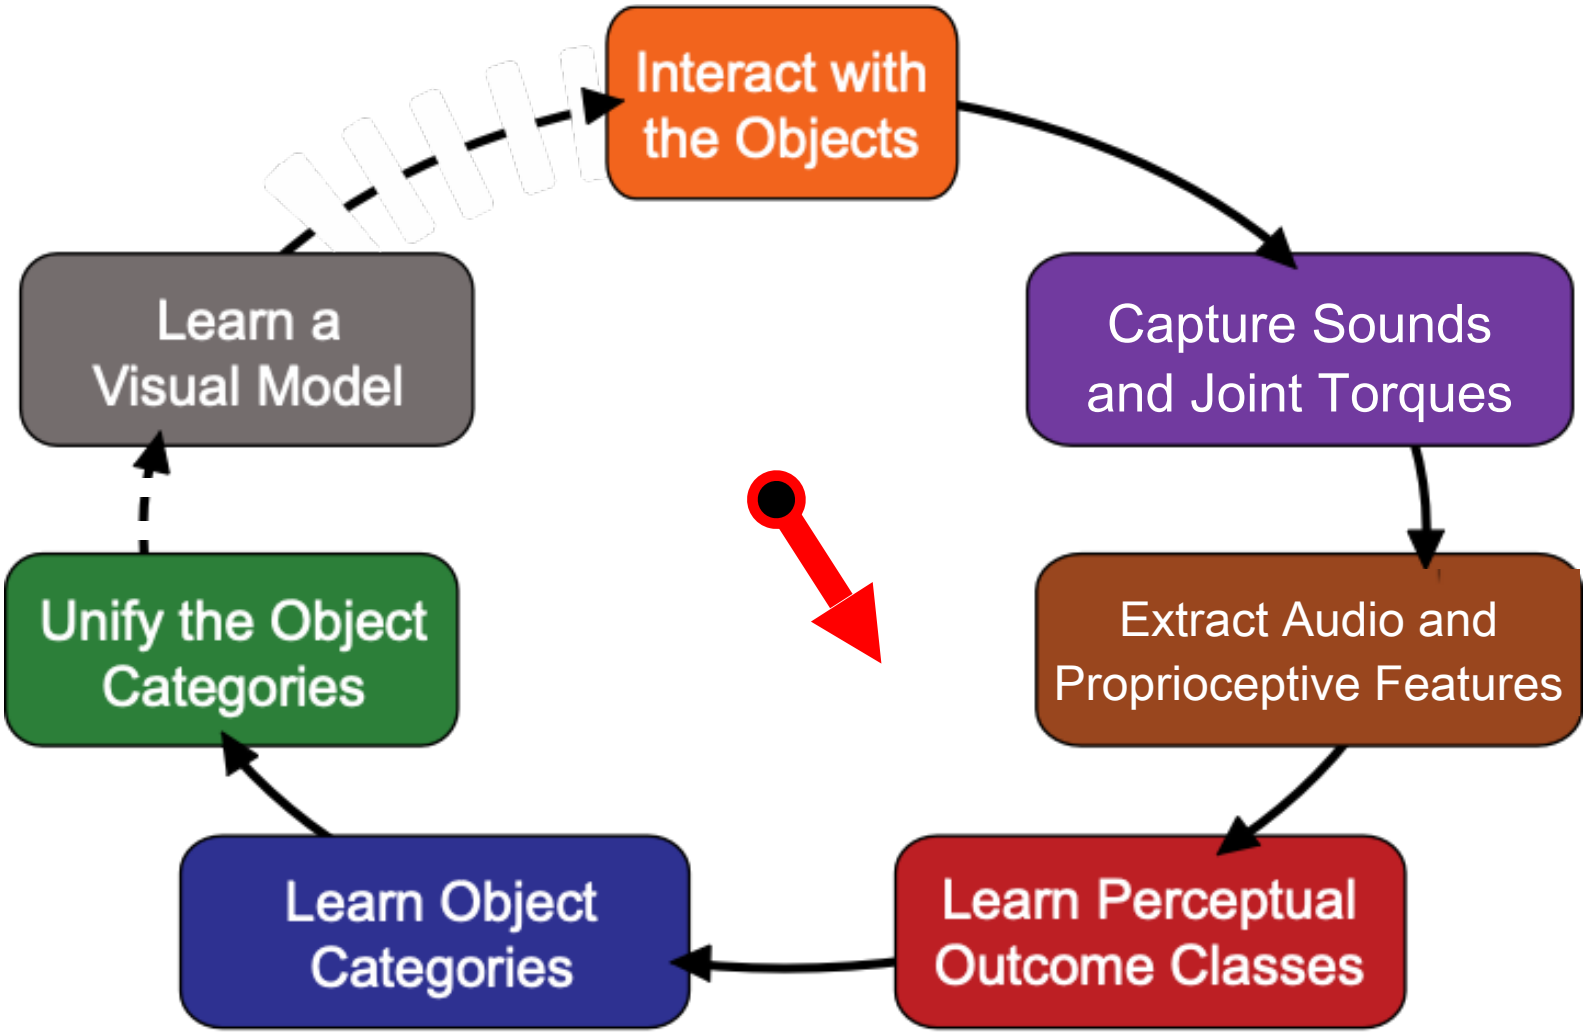

Learning Outcome Classes

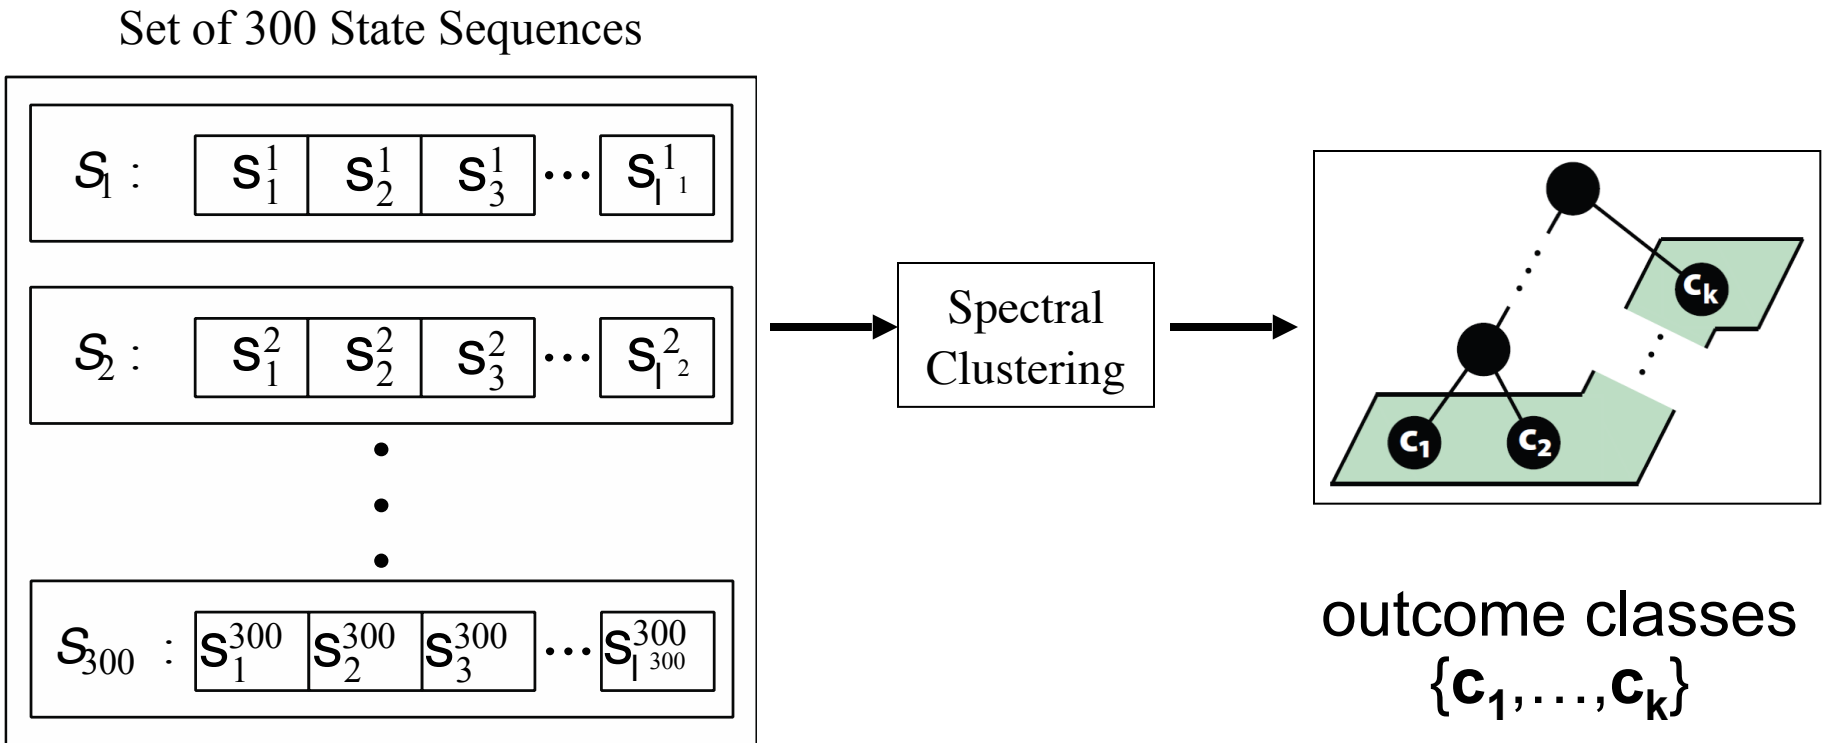

- Spectral Clustering requires a similarity matrix as input.
- Similarity function: the Needleman-Wunsch algorithm.
  (Needleman and Wunsch, 1970)

Learning Framework

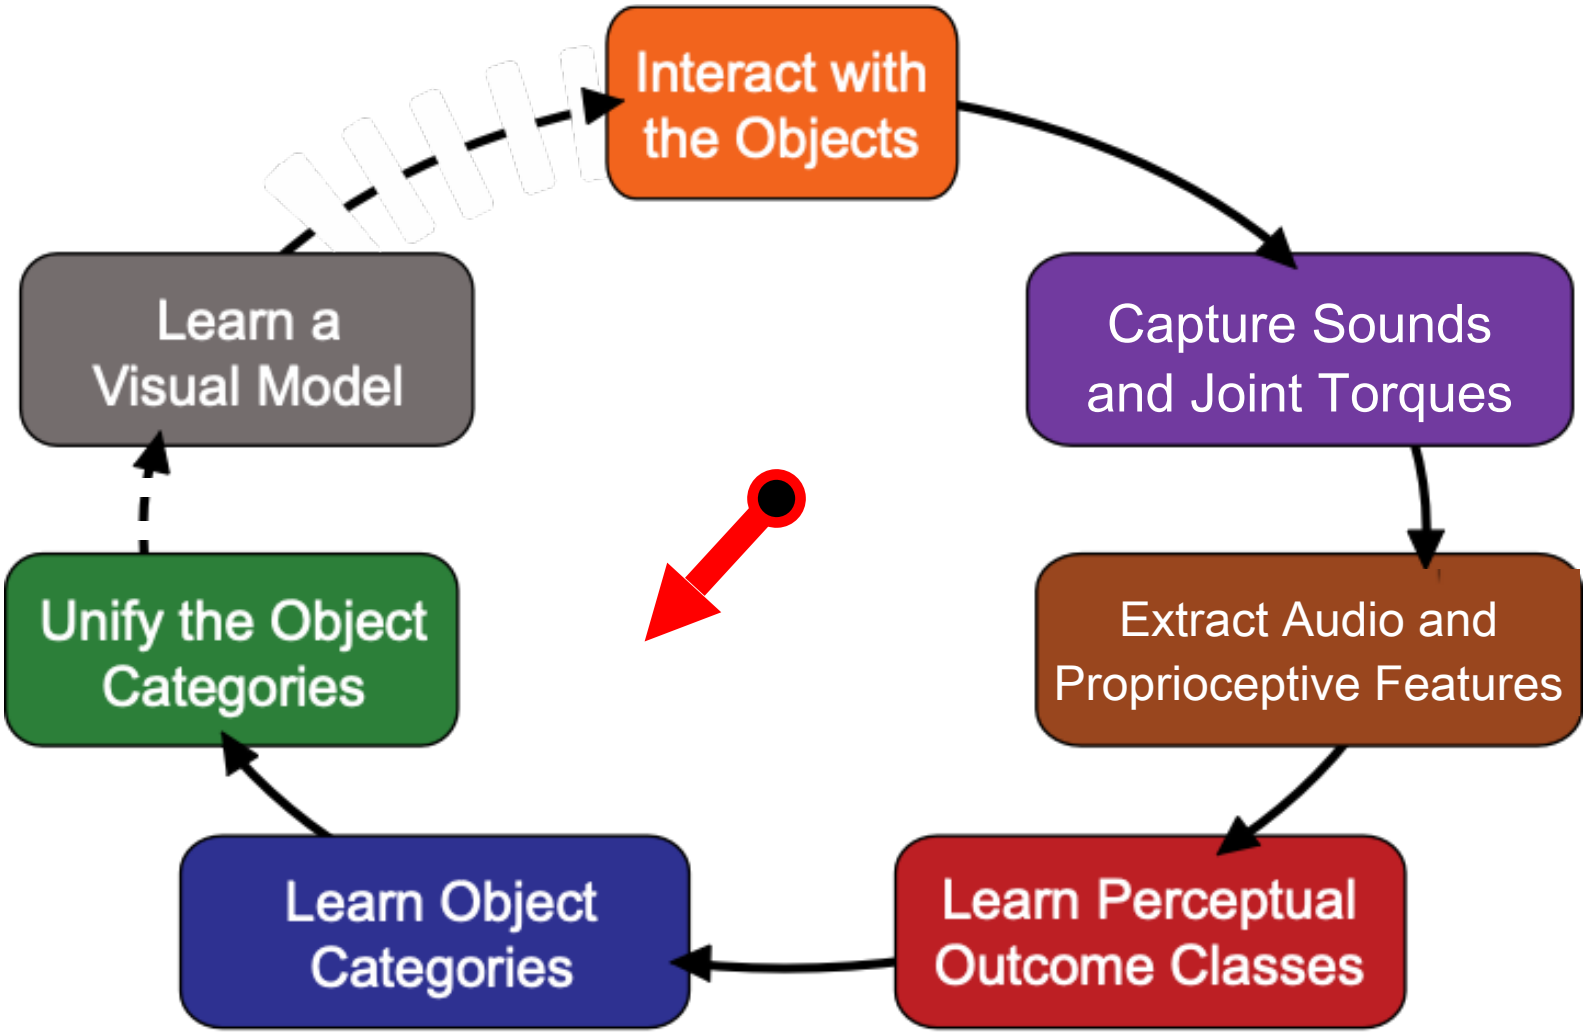

Example Outcome Classes for the In and Out Behavior

Outcome Hierarchy

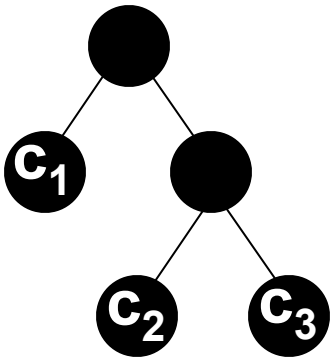

Learned Outcome Classes

- c_1 Sounds of the water filling up a tall glass
- c_2 Sounds of the water filling up a short cup
- c_3 Sounds of the water splashing against a non-container

Example Object Representation

(After 10 Repetitions of the In and Out Behavior)

Object	In and Out	Number of Outcome Class Occurrences	Object Representation
		c_1 : 0 c_2 : 0 c_3 : 10	[.0, .0, 1.0]
		c_1 : 10 c_2 : 0 c_3 : 0	[1.0, .0, .0]

Real Object Representation

(After 300 Repetitions of the In and Out Behavior)

Object

Number of Outcome Class Occurrences

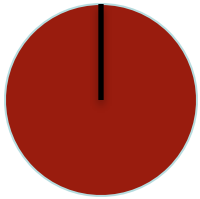

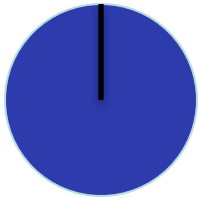

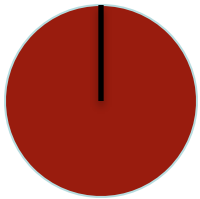

C_1

C_2

C_3

C_4

Object Categorization

X-means

(Pelleg and Moore, 2000)

Categorization Results

Cluster 1

Cluster 2

Cluster 3

Cluster 4

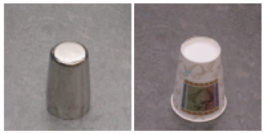

Cluster 5

Audio/In and Out

Categorization Results

Cluster 1

Cluster 2

Cluster 3

Proprioception/In and Out

Learning Framework

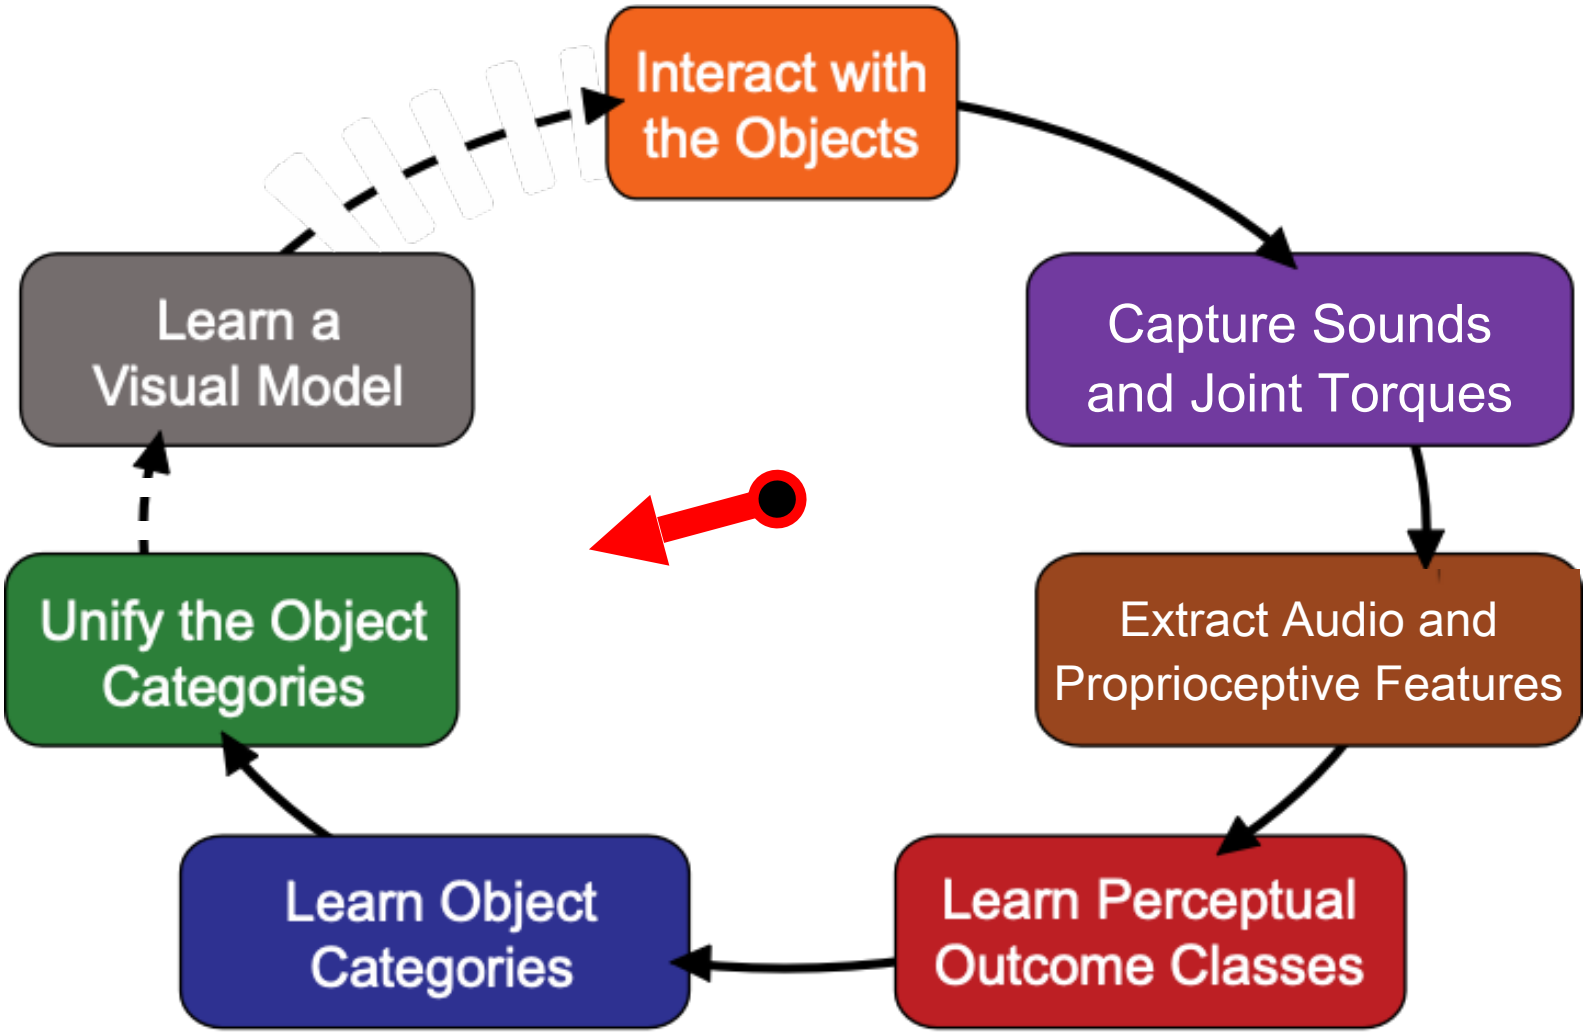

Unification Algorithm

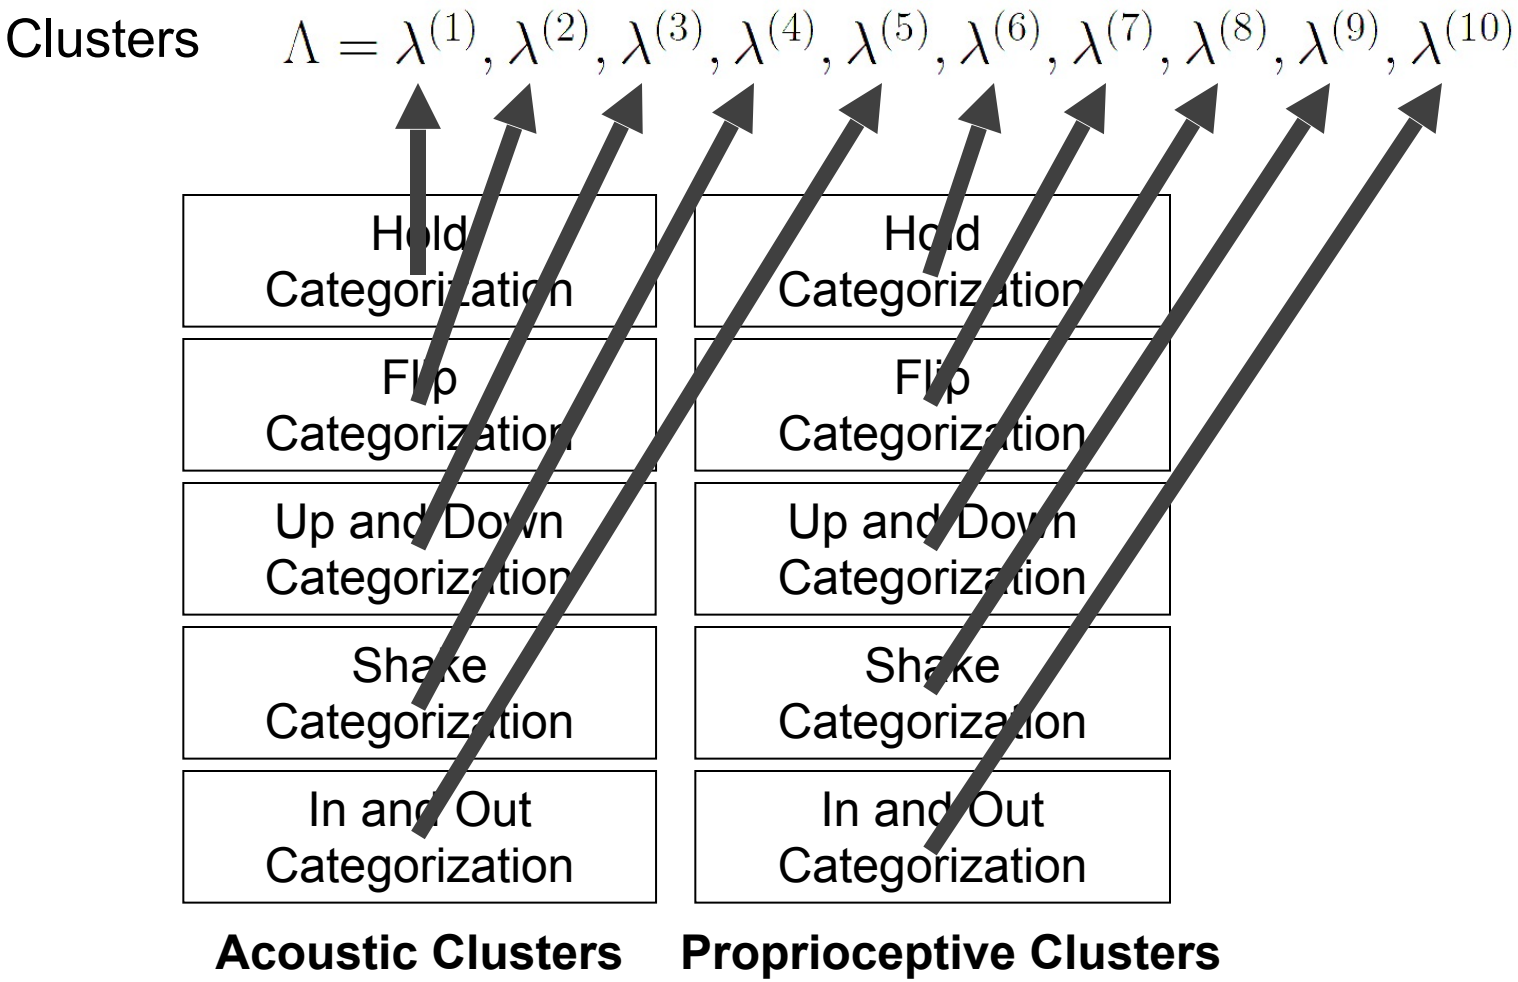

Unification Algorithm

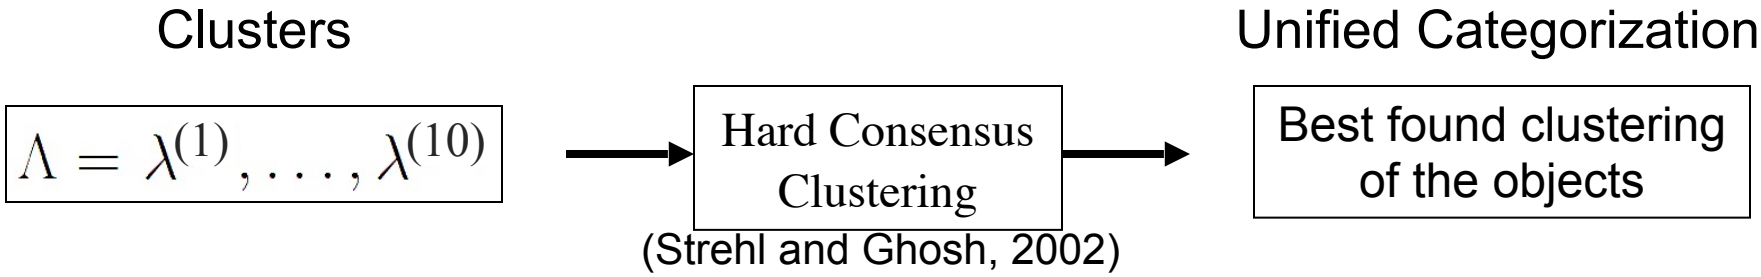

- Hard Consensus Clustering searches for a good clustering.
- The output clustering optimizes the normalized mutual information objective function.

Unified Categorization

(derived from both sound and joint torque observations)

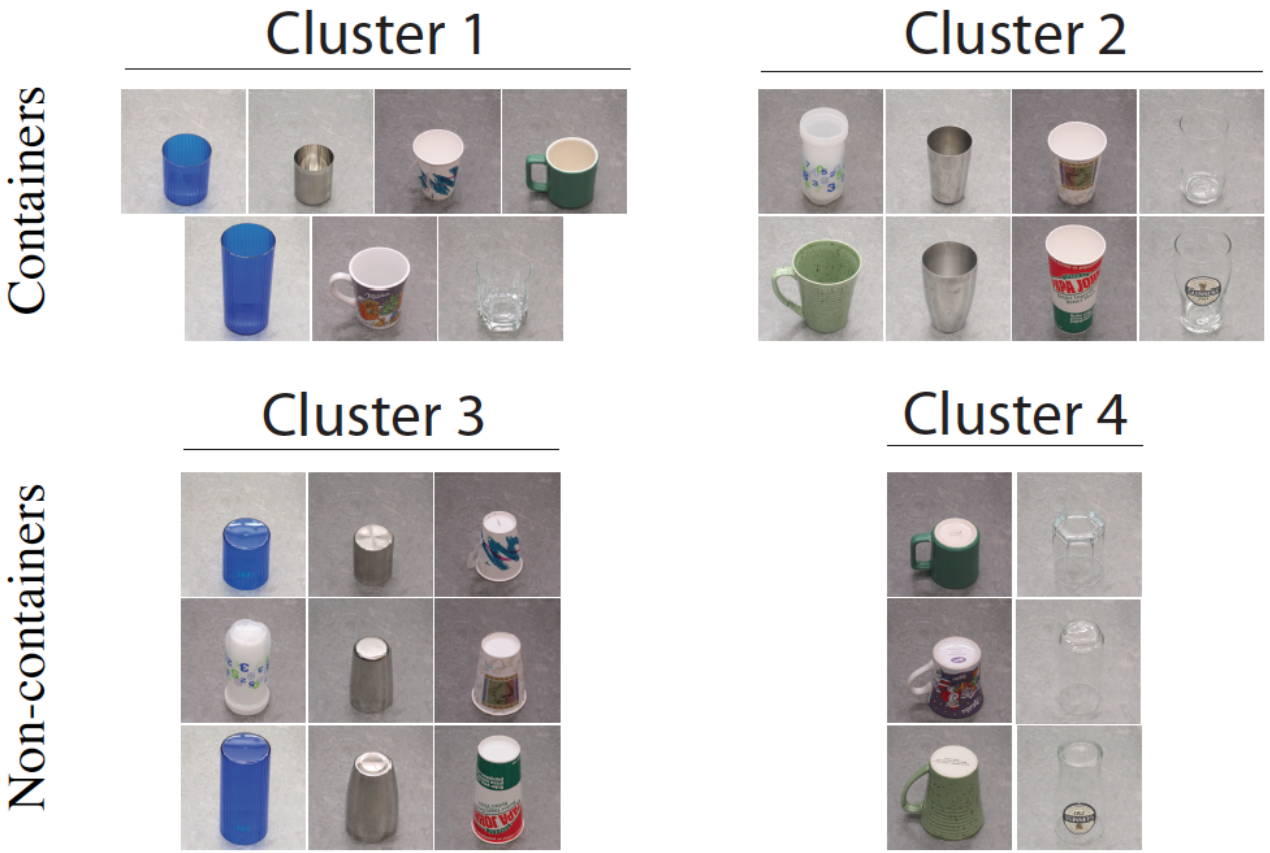

How to Evaluate the Quality of the Categorization?

- Information gain: the change in entropy from a prior state
- entropy of the prior state minus the sum of entropy computed for each cluster

UNCATEGORIZED

information gain

UNIFIED CLUSTERING

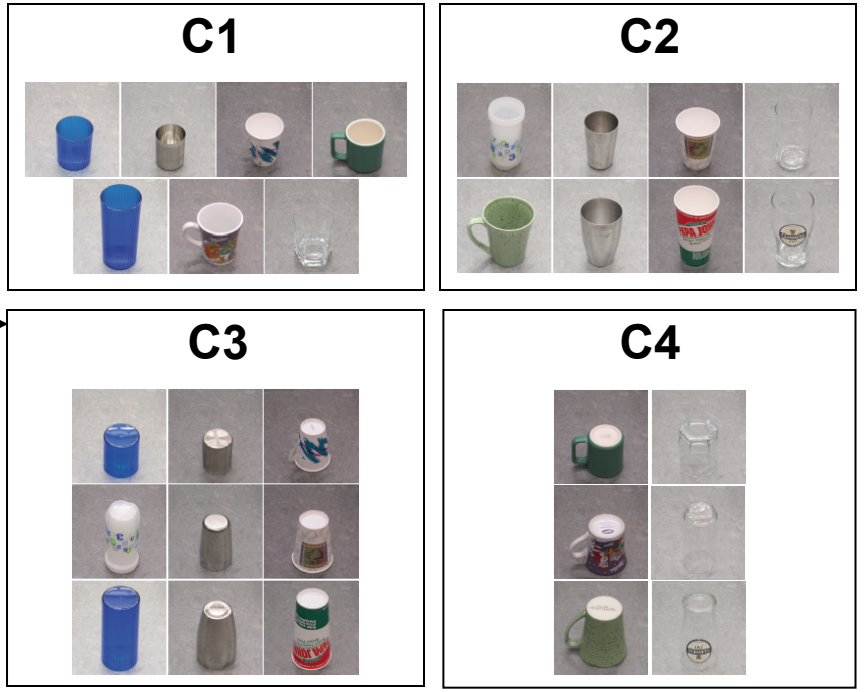

Categorization Quality

UNIFIED CLUSTERING

Cluster	Conditional Probabilities		Entropy	Information Gain
	container	non-container		
C1 	7/7	0/7		
C2 	8/8	0/8		
C3 	0/9	9/9		
C4 	0/6	6/6		

Categorization Quality

UNIFIED CLUSTERING

Cluster	Conditional Probabilities: container	Conditional Probabilities: non-container	Entropy	Information Gain
C1 	7/7	0/7	→	$-(7/7)\log_2(7/7) - (0/7)\log_2(0/7)$
C2 	8/8	0/8	→	$-(8/8)\log_2(8/8) - (0/8)\log_2(0/8)$
C3 	0/9	9/9	→	$-(0/9)\log_2(0/9) - (9/9)\log_2(9/9)$
C4 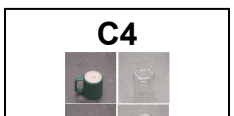	0/6	6/6	→	$-(0/6)\log_2(0/6) - (6/6)\log_2(6/6)$

Categorization Quality

UNIFIED CLUSTERING

Cluster	Conditional Probabilities		Entropy	Information Gain
	container	non-container		
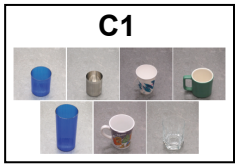	7/7	0/7	→ 0.0	
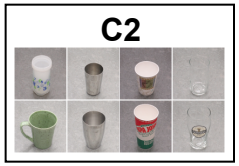	8/8	0/8	→ 0.0	
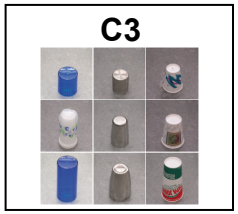	0/9	9/9	→ 0.0	
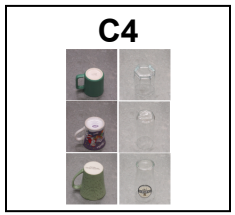	0/6	6/6	→ 0.0	

Categorization Quality

UNIFIED CLUSTERING

Cluster	Conditional Probabilities: container	Conditional Probabilities: non-container	Entropy	Information Gain
C1 	7/7	0/7	0.0	– (7/30 * 0.0)
C2 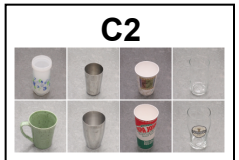	8/8	0/8	0.0	– (8/30 * 0.0)
C3 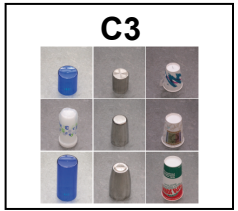	0/9	9/9	0.0	– (9/30 * 0.0)
C4 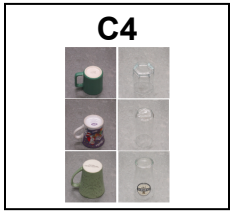	0/6	6/6	0.0	– (6/30 * 0.0)

1 – (7/30 * 0.0) – (8/30 * 0.0) – (9/30 * 0.0) – (6/30 * 0.0)

Categorization Quality

UNIFIED CLUSTERING

Cluster	Conditional Probabilities: container	Conditional Probabilities: non-container	Entropy	Information Gain
C1 	7/7	0/7	0.0	– (7/30 * 0.0)
C2 	8/8	0/8	0.0	– (8/30 * 0.0)
C3 	0/9	9/9	0.0	– (9/30 * 0.0)
C4 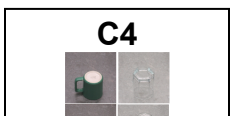	0/6	6/6	0.0	– (6/30 * 0.0)

1.0

Information Gained with Respect to Human Labels

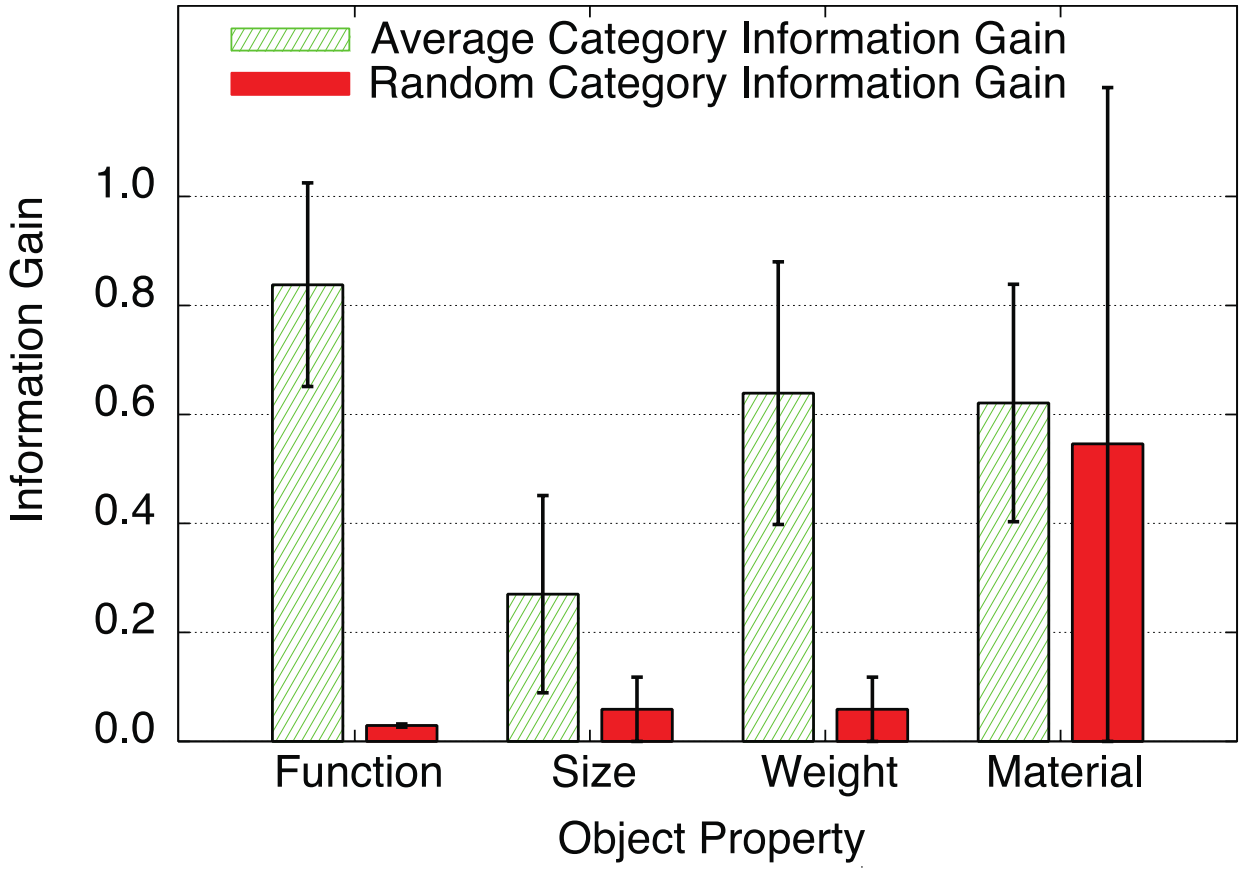

Conclusion

- Object category learning is possible while interacting with objects in a sink.
- The sounds and the sensations of water flowing into a cup is one embodiment of that object.
- Sound and proprioception might be able to bootstrap water manipulation research.

Acknowledgements

NSF

Questions?

Material Properties of the Objects

X-means (Pelleg and Moore, 2000)

K-means Result

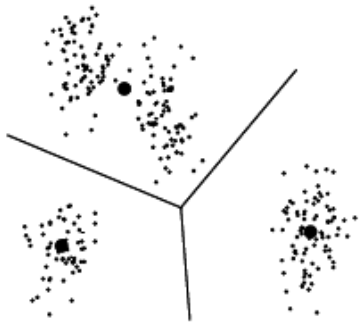

Split the Centroids

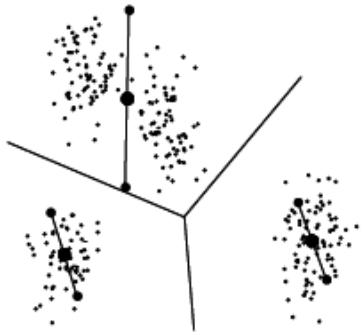

Run K-means locally

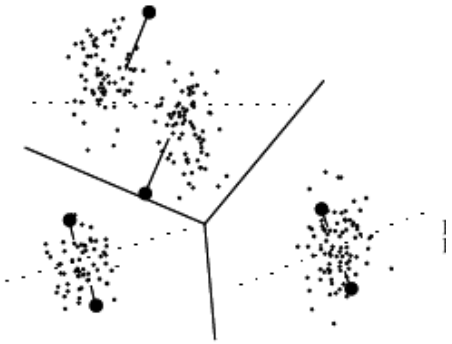

Evaluate the New Centroids

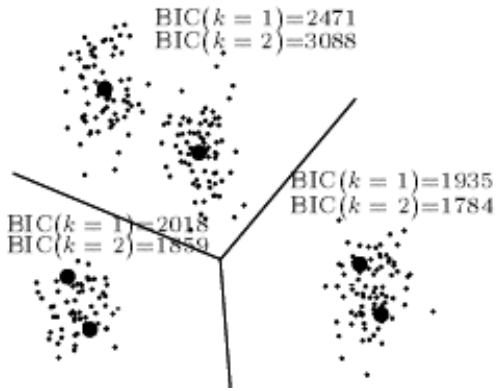

New Structure

Needleman-Wunsch (1970)

Initialize the Matrix

		C	O	E	L	A	C	A	N	T	H
	0	-1	-2	-3	-4	-5	-6	-7	-8	-9	-10
P	-1										
E	-2										
L	-3										
I	-4										
C	-5										
A	-6										
N	-7										

Fill the Matrix

		C	O	E	L	A	C	A	N	T	H
	0	-1	-2	-3	-4	-5	-6	-7	-8	-9	-10
P	-1	**-1**	-2	-3	-4	-5	-6	-7	-8	-9	-10
E	-2	-2	-2	-1	-0	-3	-4	-5	-6	-7	-8
L	-3	-3	-3	-2	-2	-1	-2	-3	-4	-5	-6
I	-4	-4	-4	-3	-1	-1	-2	-1	-4	-5	-6
C	-5	-3	-4	-4	-2	-2	-0	-1	-2	-3	-4
A	-6	-4	-4	-5	-3	-1	-1	-1	-0	-1	-2
N	-7	-5	-5	-5	-4	-2	-2	-0	-2	-1	-0

What Sensory Modalities to Use?

Vision

Audio

Proprioception

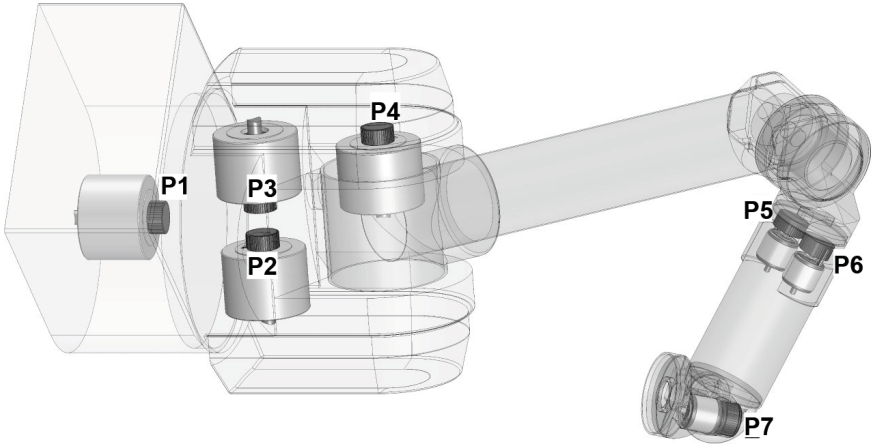

Learning Framework

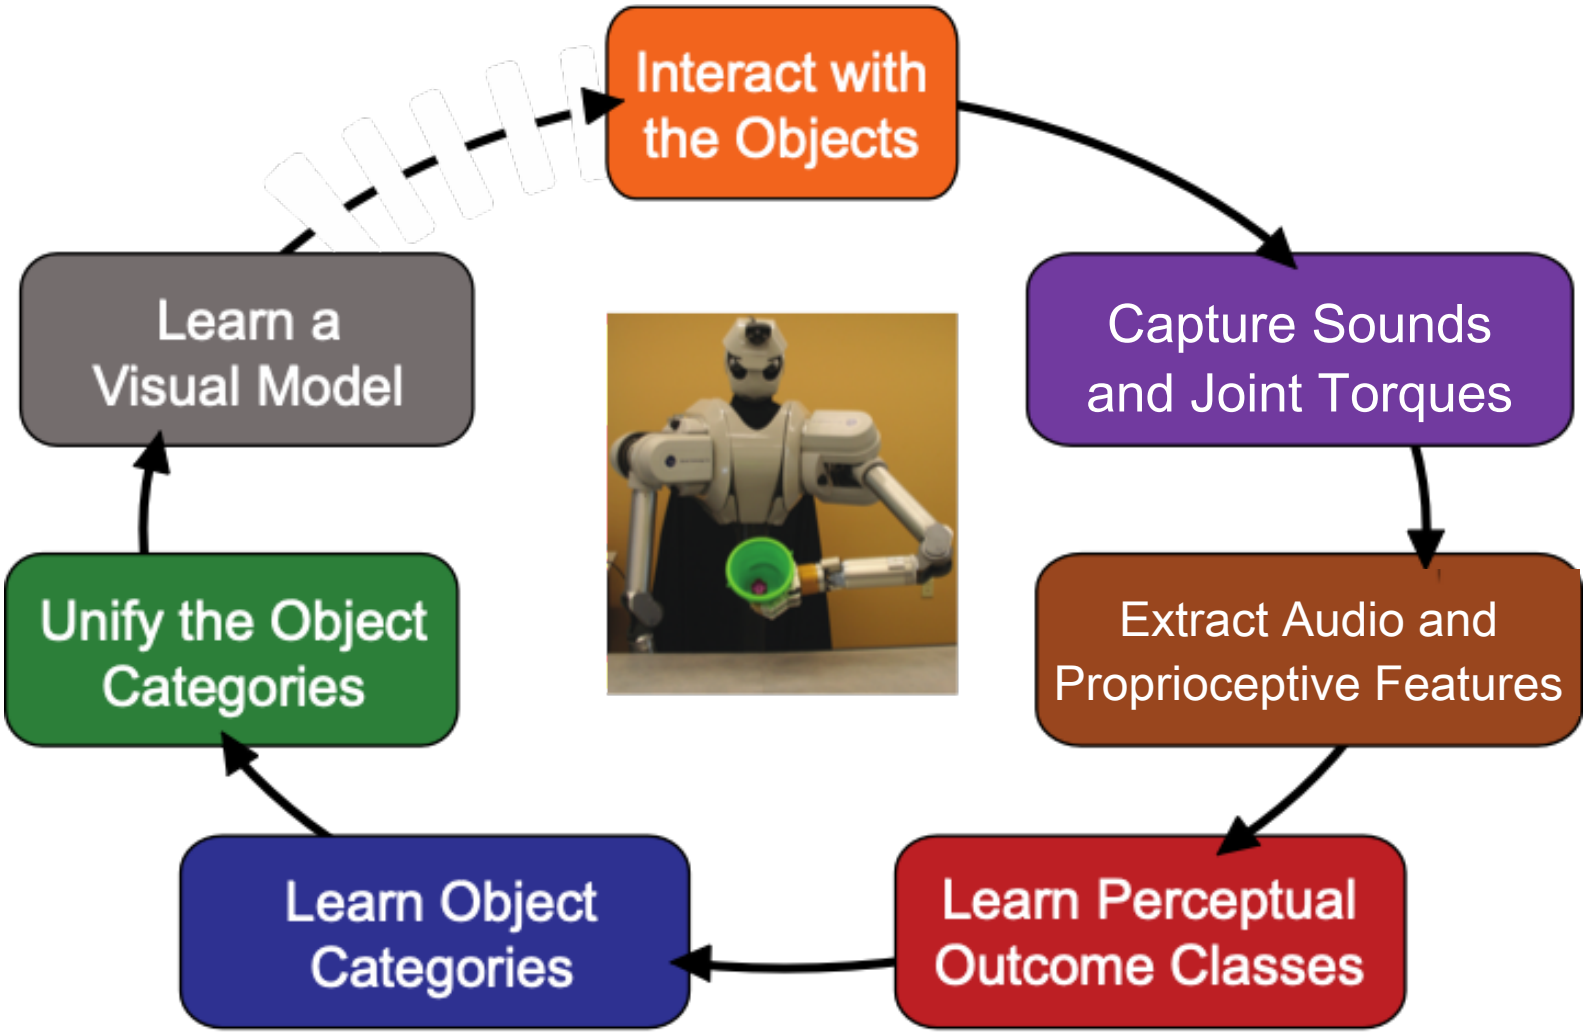

Learning Framework